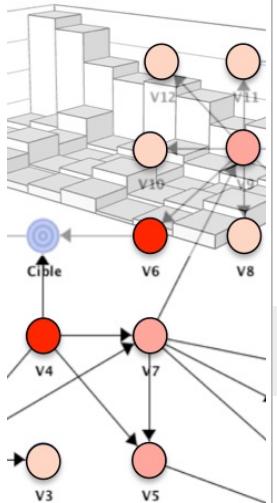


$$+ \sqrt{\frac{c_e \ln(T_F)}{t_{F,a}}} \min\left(\frac{1}{4}, \hat{\mu}_{F,a} + \sqrt{\frac{2 \ln(T_F)}{t_{F,a}}}\right)$$



```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.  
        .
```

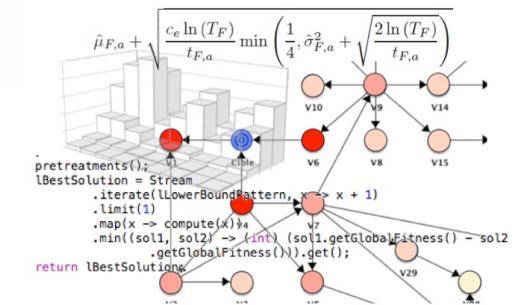
$$= \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

DEEP LEARNING

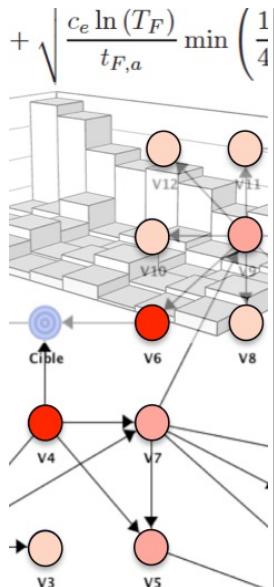
Transfer Learning

Stéphane BONNEVAY

Polytech Lyon
Laboratoire ERIC
stephane.bonnevay@univ-lyon1.fr

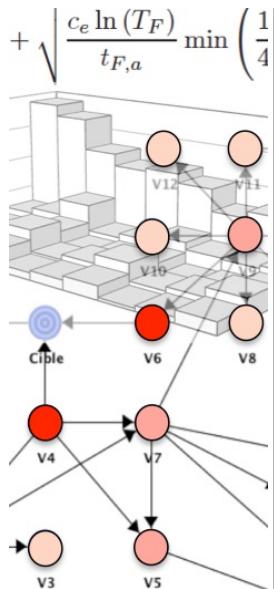


```
.pretreatments();  
lBestSolution = Stream  
    .iterate(lLowerBoundPattern,  
    .limit(1),  
    .map(x -> compute(x)).  
    .min((sol1, sol2) -> (int) (sol1.getGlobalFitness() - sol2  
    .getGlobalFitness())).get();  
return lBestSolutions;
```



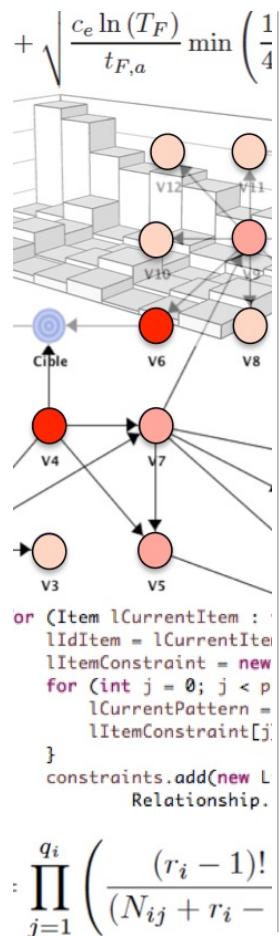
```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.  
        ...);  
}
```

$$= \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

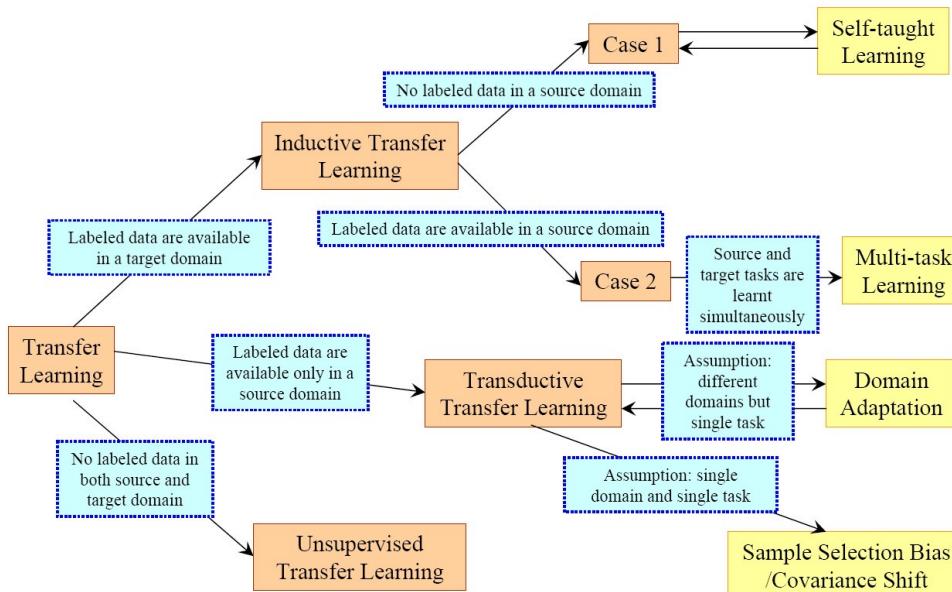


```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.  
        ...);  
}
```

$$= \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$



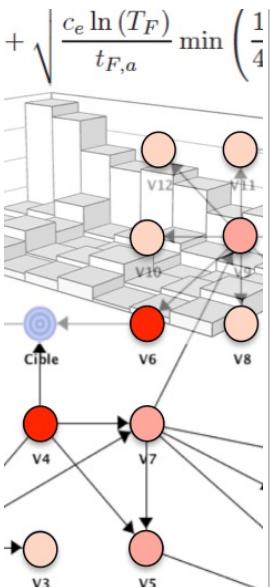
Le **Transfer Learning** consiste à réutiliser l'expérience acquise au terme d'un apprentissage dans le cadre d'un domaine ou d'une tâche.



Pan and Q. Yang. « A survey on Transfer Learning », IEEE Transactions on Knowledge and Data Engineering, Vol.22(10), October, 2010.

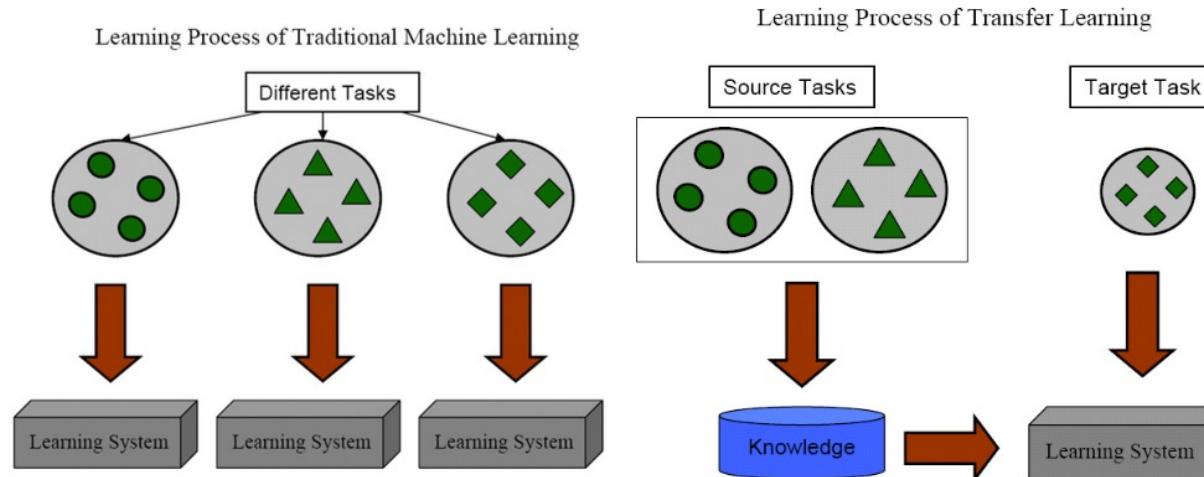
Idée :

adapter un modèle existant, destiné à une certaine tâche et à un certain domaine
à un autre domaine pour réaliser une autre tâche

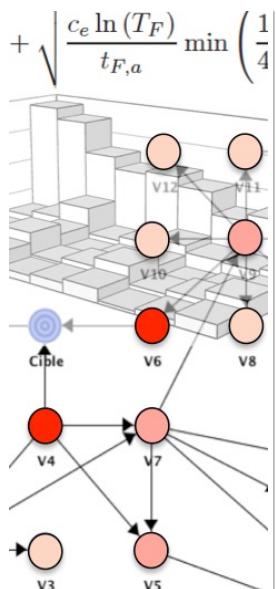


```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.
```

$$= \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$



Pan and Q. Yang. « A survey on Transfer Learning », IEEE Transactions on Knowledge and Data Engineering, Vol.22(10), October, 2010.



```

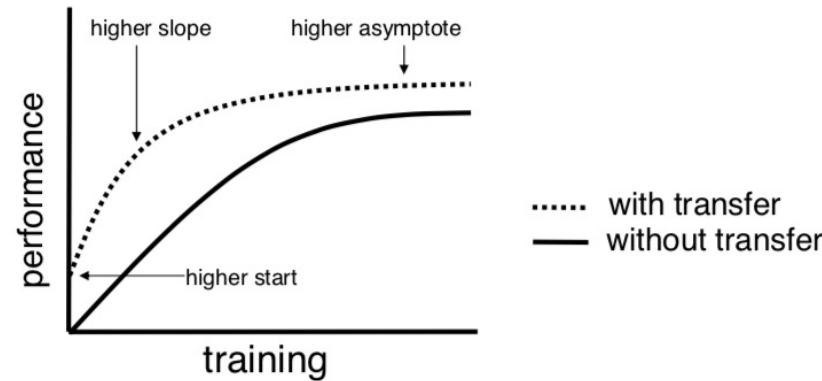
or (Item lCurrentItem :
    lIdItem = lCurrentItem
    lItemConstraint = new
    for (int j = 0; j < p
        lCurrentPattern =
        lItemConstraint[j];
    }
    constraints.add(new L
        Relationship.

```

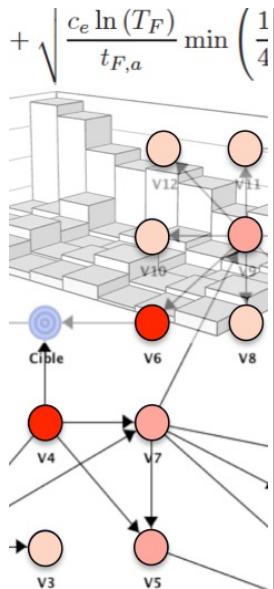
$$= \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

Avantages :

- amélioration de la qualité de l'apprentissage
- adaptation de « gros » modèles nécessitant des moyens techniques hors de portée pour son propre problème
- réduction du temps de développement d'un modèle
- performance finale accrue



Torrey, L. and Shavlik, J. "Transfer Learning. In: Handbook of Research on Machine Learning Applications and Trends: Algorithms, Methods, and Techniques", IGI Global, Hershey, 242-264, 2010.

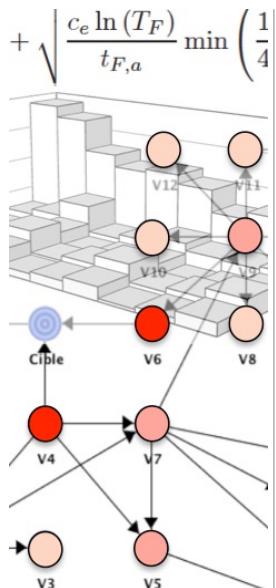


```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.  
        .
```

$$: \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

1. Introduction
2. Classification d'images avec
 - un réseau multi-couches
 - un réseau de neurones convolutifs
 - un réseau pré-entraîné (transfer learning)
3. Détection d'objets dans une image
4. Classification de textes

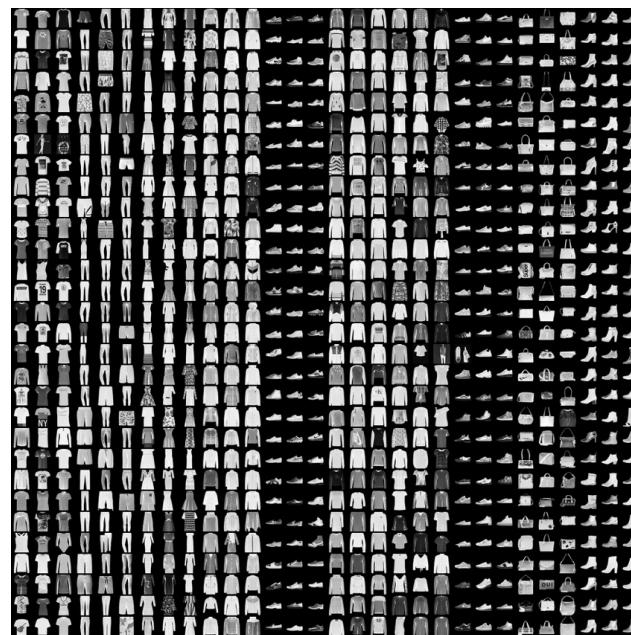
Deep Learning - Transfer Learning



```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.  
        .
```

$$: \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

Exemple : jeu de données « Fashion MNIST »



Dataset : <https://github.com/zalandoresearch/fashion-mnist>

Stéphane BONNEVAY – Polytech Lyon

Classification d'images - Réseaux multicouches

70 000 images en niveau de gris de taille 28x28

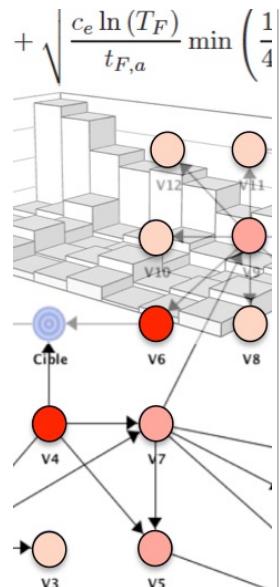
valeurs de 0 à 255

60 000 pour l'apprentissage

10 000 pour le test

```
from tensorflow.keras.datasets import fashion_mnist  
(x_train, y_train), (x_test, y_test) = fashion_mnist.load_data()
```

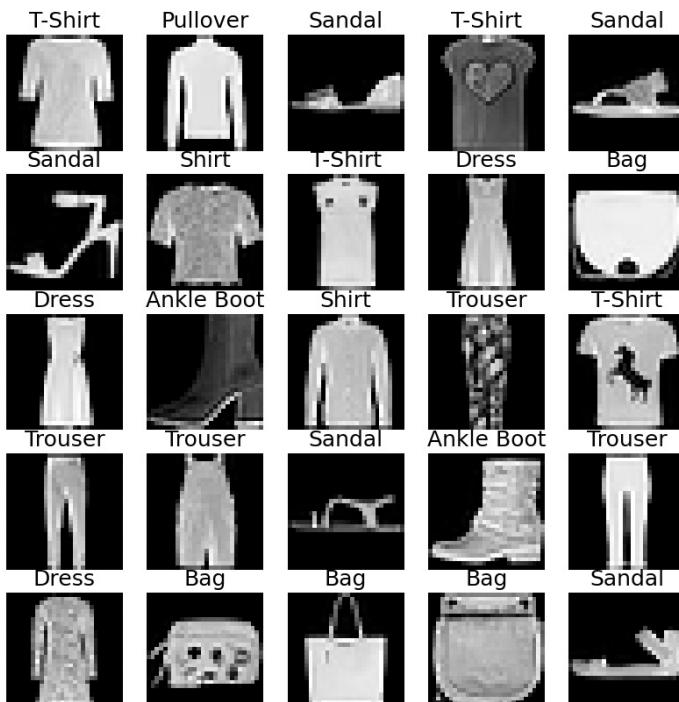
10 catégories



```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.
```

$$: \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

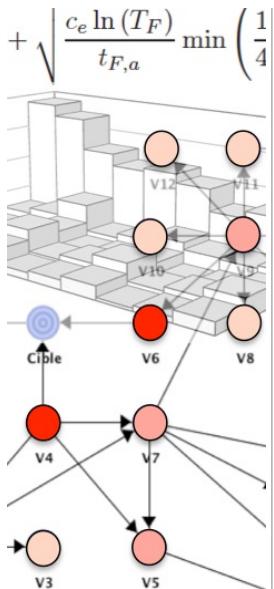
Exemple : jeu de données « Fashion MNIST »



10 catégories :

- 0: T-shirt
- 1: Trouser
- 2: Pullover
- 3: Dress
- 4: Coat
- 5: Sandal
- 6: Shirt
- 7: Sneaker
- 8: Bag
- 9: Ankle boot

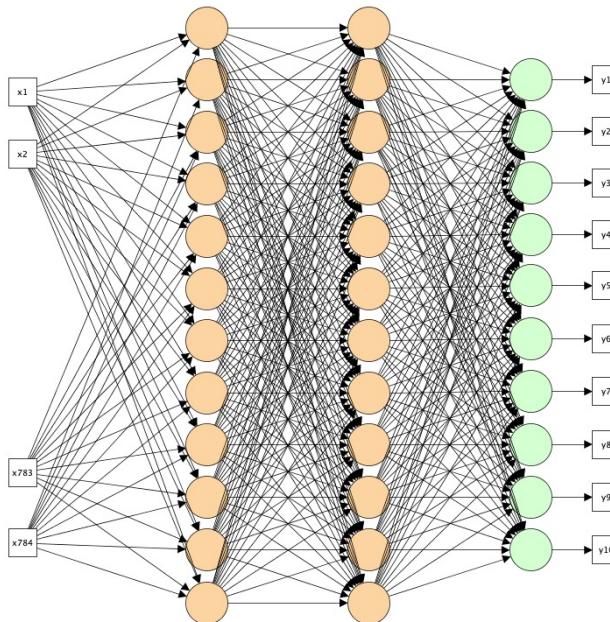
Deep Learning - Transfer Learning



```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.  
    )
```

$$= \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

Stéphane BONNEVAY – Polytech Lyon



Paramètres conseillés :

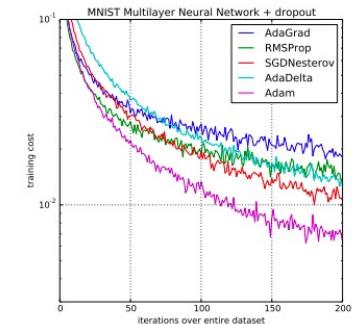
Learning rate = 0.001
Beta = 0.9
Beta2 = 0.999
Epsilon = 1e-08

Classification d'images - Réseaux multicouches

Il est possible de combiner différents types de fonctions d'activation (à ce jour, ReLu est la plus utilisée).

Il existe plusieurs algorithmes d'apprentissage :

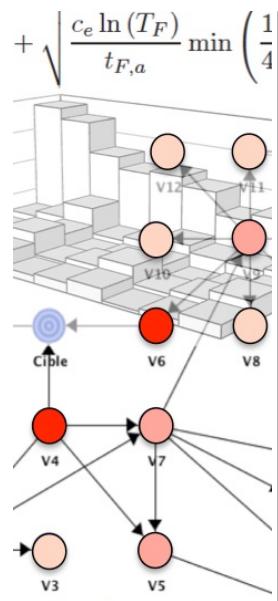
- AdaGrad
- Adam ← le plus utilisé
- RMSProp
- AdaDelta
- ...



in « A Method for Stochastic Optimization », 2015

Deep Learning - Transfer Learning

Classification d'images - Réseaux multicouches



```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.
```

$$: \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

Réseau de neurones : 1 couche cachée, 10 neurones



Matrice de taille

28x28

→ Vecteur d'entrée de taille 784

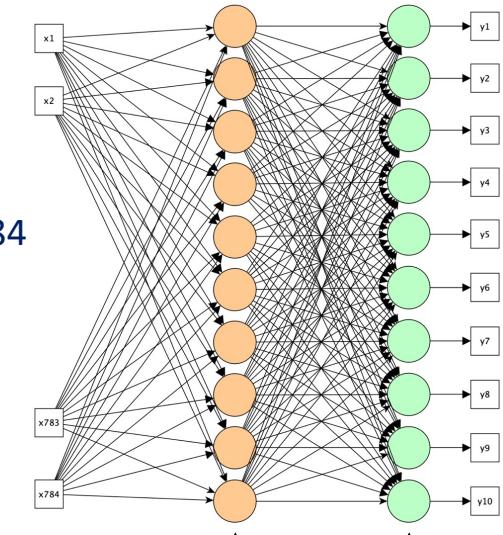
$54\ 000 \times 28 \times 28 \rightarrow 54\ 000 \times 784$

$6\ 000 \times 28 \times 28 \rightarrow 6\ 000 \times 784$

$10\ 000 \times 28 \times 28 \rightarrow 10\ 000 \times 784$

```
x_train = x_train.reshape(x_train.shape[0], -1) / 255.0  
x_test = x_test.reshape(x_test.shape[0], -1) / 255.0  
y_train = to_categorical(y_train)  
y_test = to_categorical(y_test)
```

Stéphane BONNEVAY – Polytech Lyon

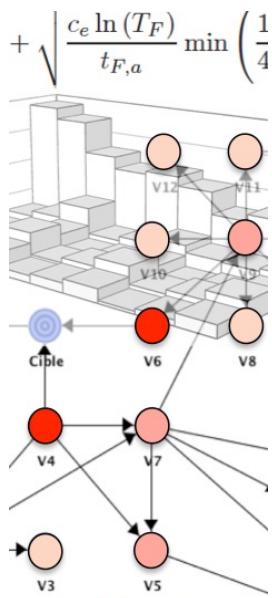


Choix arbitraire de 10 neurones
sur la couche cachée

Couche de sortie de taille 10
une par catégorie

Deep Learning - Transfer Learning

Classification d'images - Réseaux multicouches



```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.  
    )
```

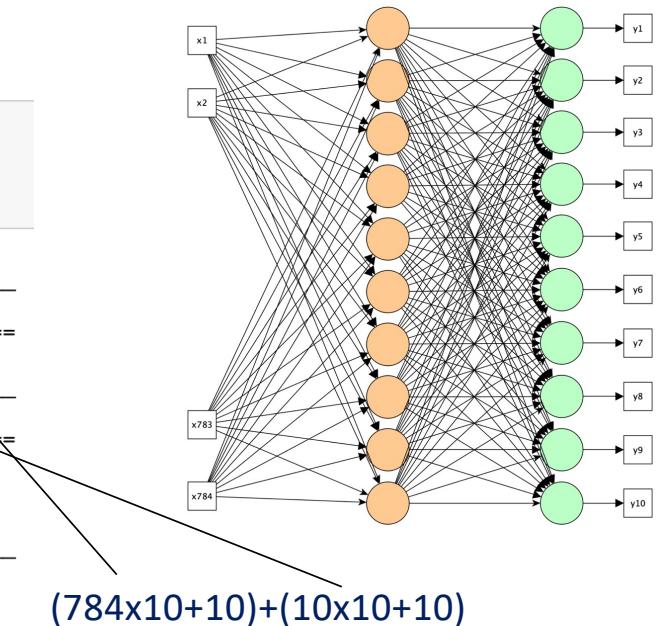
$$= \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

Réseau de neurones : 1 couche cachée, 10 neurones

```
model = Sequential()  
model.add(Dense(10, input_dim=784, activation='relu'))  
model.add(Dense(10, activation='softmax'))  
model.summary()
```

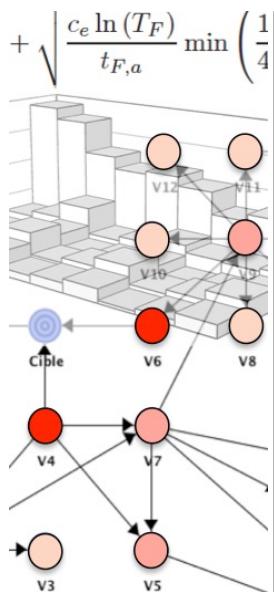
Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 10)	7850
dense_1 (Dense)	(None, 10)	110
Total params:	7,960	
Trainable params:	7,960	
Non-trainable params:	0	



Fonctions d'activation

- « Relu » sur les couches cachées
- « Softmax » si plus de 2 catégories sur la couche de sortie
- « Sigmoid » si 2 catégories sur la couche de sortie



```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.
```

$$: \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

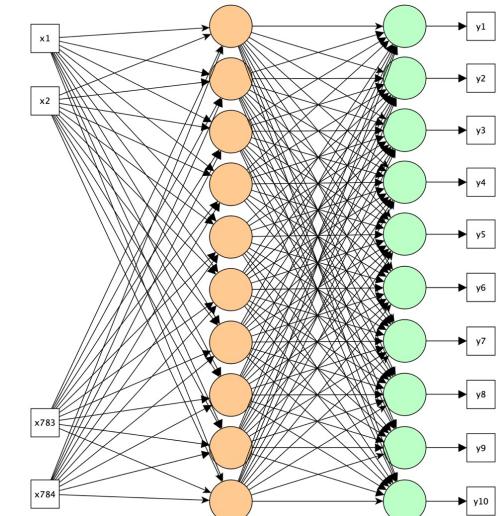
Réseau de neurones : 1 couche cachée, 10 neurones

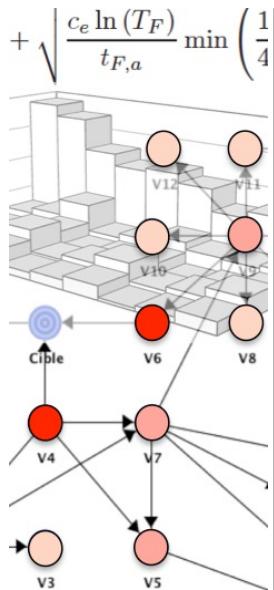
```
model.compile(loss='categorical_crossentropy',
              optimizer='adam',
              metrics=['accuracy'])
```

Paramètres

- « categorical_crossentropy » pour du multi-catégories
- « binary_crossentropy » pour 2 catégories
- « adam » ou « rmsprop » sont les 2 meilleures
- « accuracy » pour la métrique

$$\text{accuracy} = \frac{\text{nombre de bonnes prédictions}}{\text{nombre total de prédictions}}$$





```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.  
    )
```

$$= \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

Réseau de neurones : 1 couche cachée, 10 neurones

```
model.fit(x_train, y_train,  
          epochs=10,  
          batch_size=32,  
          validation_split=0.1)
```

Paramètres

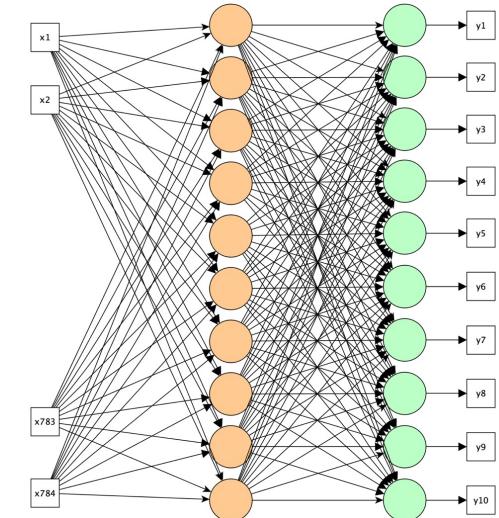
10 itérations (epochs)

1 epoch = passage de l'ensemble des données dans le réseau

Batch_size = nombre d'enregistrements qui passent dans le réseau avant la mise à jour des poids

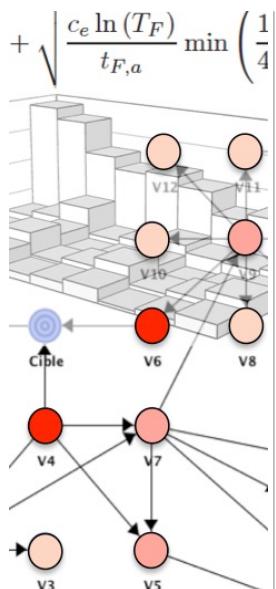
10% des données d'apprentissage pour la validation après chaque epoch

54 000 images pour l'apprentissage
6 000 images pour la validation



Deep Learning - Transfer Learning

Classification d'images - Réseaux multicouches



```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.
```

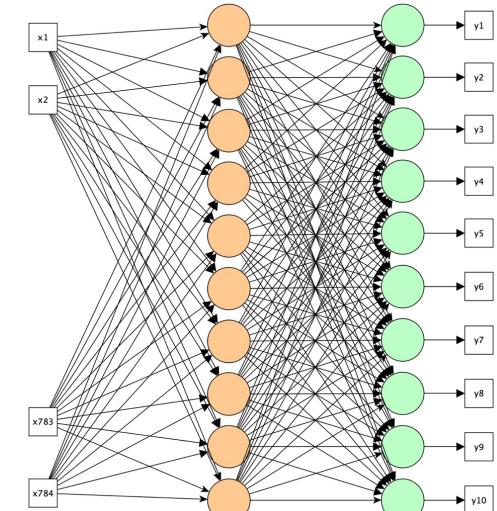
$$: \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

Stéphane BONNEVAY – Polytech Lyon

Réseau de neurones : 1 couche cachée, 10 neurones

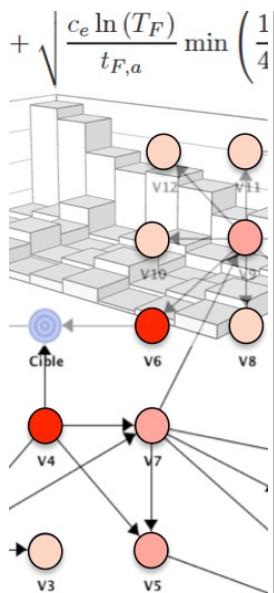
```
Train on 54000 samples, validate on 6000 samples
Epoch 1/10
54000/54000 [=====] - 3s 49us/sample - loss: 0.7834 - accuracy: 0.7354 - val_loss: 0.5066 -
val_accuracy: 0.8220
Epoch 2/10
54000/54000 [=====] - 2s 32us/sample - loss: 0.4820 - accuracy: 0.8331 - val_loss: 0.4505 -
val_accuracy: 0.8402
Epoch 3/10
54000/54000 [=====] - 2s 33us/sample - loss: 0.4457 - accuracy: 0.8447 - val_loss: 0.4391 -
val_accuracy: 0.8472
Epoch 4/10
54000/54000 [=====] - 2s 35us/sample - loss: 0.4264 - accuracy: 0.8506 - val_loss: 0.4296 -
val_accuracy: 0.8453
Epoch 5/10
54000/54000 [=====] - 2s 34us/sample - loss: 0.4153 - accuracy: 0.8546 - val_loss: 0.4217 -
val_accuracy: 0.8505
Epoch 6/10
54000/54000 [=====] - 2s 34us/sample - loss: 0.4067 - accuracy: 0.8584 - val_loss: 0.4216 -
val_accuracy: 0.8543
Epoch 7/10
54000/54000 [=====] - 2s 36us/sample - loss: 0.4011 - accuracy: 0.8601 - val_loss: 0.4137 -
val_accuracy: 0.8527
Epoch 8/10
54000/54000 [=====] - 2s 33us/sample - loss: 0.3961 - accuracy: 0.8615 - val_loss: 0.4396 -
val_accuracy: 0.8430
Epoch 9/10
54000/54000 [=====] - 2s 36us/sample - loss: 0.3902 - accuracy: 0.8622 - val_loss: 0.4120 -
val_accuracy: 0.8542
Epoch 10/10
54000/54000 [=====] - 2s 33us/sample - loss: 0.3863 - accuracy: 0.8637 - val_loss: 0.4095 -
val_accuracy: 0.8555
```

85,55% de précision sur les données de validation



Deep Learning - Transfer Learning

Classification d'images - Réseaux multicouches



```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p;  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.
```

$$= \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

Réseau de neurones : 1 couche cachée, 10 neurones

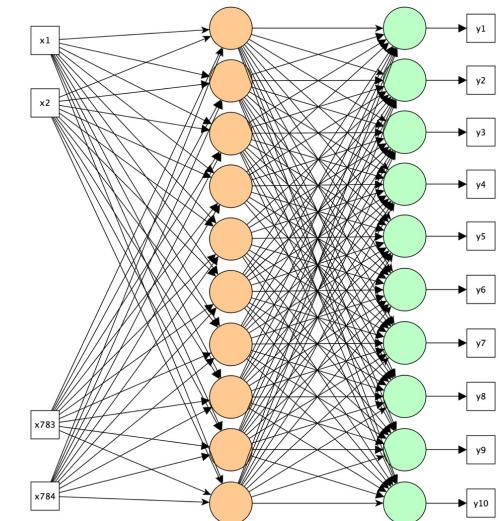
```
_ , test_acc = model.evaluate(x_test, y_test)  
print(test_acc)
```

- 0s 18us/sample - loss: 0.3917 - accuracy: 0.8460

84,60% de précision sur les données de test

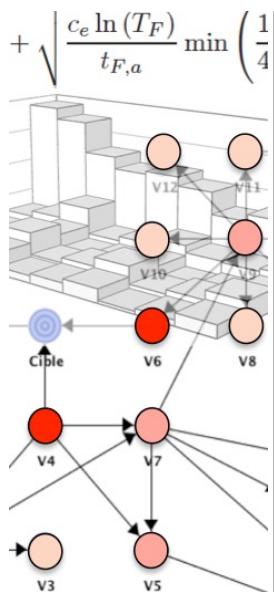


Le réseau de neurones classe correctement 8460 images sur les 10000 images des données de test



Deep Learning - Transfer Learning

Classification d'images - Réseaux multicouches



```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.  
        ...)
```

$$= \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

Réseau de neurones : 1 couche cachée, 50 neurones

```
model2 = Sequential()  
model2.add(Dense(50, input_dim=784, activation='relu'))  
model2.add(Dense(10, activation='softmax'))  
model2.summary()
```

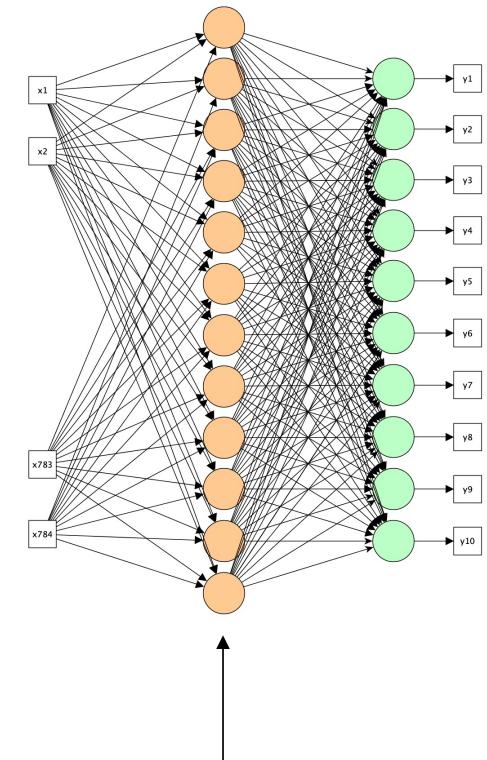
Model: "sequential_2"

Layer (type)	Output Shape	Param #
<hr/>		
dense_4 (Dense)	(None, 50)	39250
dense_5 (Dense)	(None, 10)	510
<hr/>		
Total params: 39,760		
Trainable params: 39,760		
Non-trainable params: 0		

88,05% de précision sur les données de validation

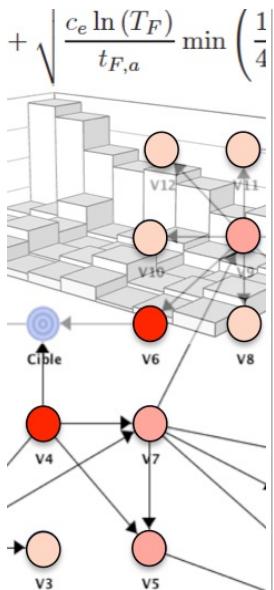
87,43% de précision sur les données de test

Stéphane BONNEVAY – Polytech Lyon



50 neurones sur la couche cachée
au lieu de 10

Deep Learning - Transfer Learning



```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.  
    )
```

$$= \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

Stéphane BONNEVAY – Polytech Lyon

Classification d'images - Réseaux multicouches

Réseau de neurones : 2 couches cachées, 50 neurones

```
model3 = Sequential()  
model3.add(Dense(50, input_dim=784, activation='relu'))  
model3.add(Dense(50, activation='relu'))  
model3.add(Dense(10, activation='softmax'))  
model3.summary()
```

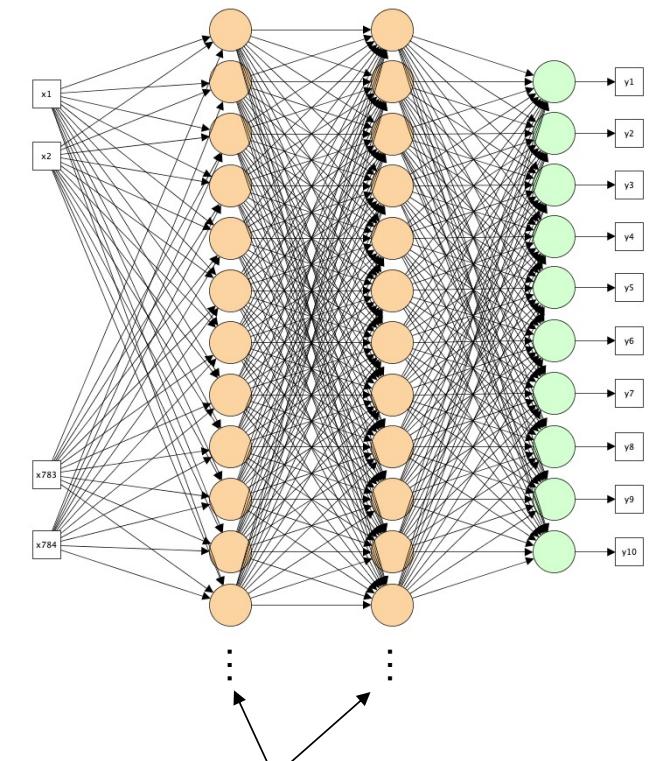
Model: "sequential_3"

Layer (type)	Output Shape	Param #
dense_6 (Dense)	(None, 50)	39250
dense_7 (Dense)	(None, 50)	2550
dense_8 (Dense)	(None, 10)	510
Total params: 42,310		
Trainable params: 42,310		
Non-trainable params: 0		

88,25% de précision sur les données de validation

87,64% de précision sur les données de test

Résultats quasi identiques qu'avec une seule couche



2 couches cachées de
50 neurones chacune

$$+ \sqrt{\frac{c_e \ln(T_F)}{t_{F,a}}} \min \left(\frac{1}{4}, \dots \right)$$

```

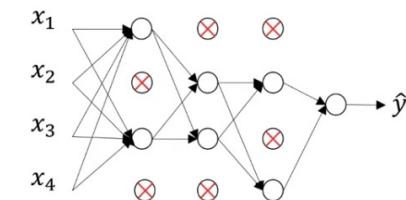
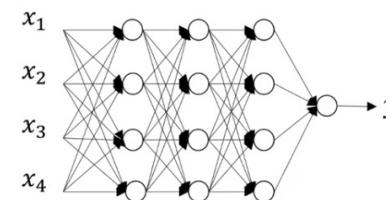
or (Item lCurrentItem : 
    lIdItem = lCurrentItem
    lItemConstraint = new
    for (int j = 0; j < p
        lCurrentPattern =
        lItemConstraint[j];
    }
    constraints.add(new L
        Relationship.
    )
    q_i
    \prod_{j=1}^{q_i} \left( \frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)
  
```

Si les résultats montrent que le modèle n'est pas bon en généralisation (mauvaise précision sur les données de test), il est possible d'appliquer quelques techniques de **régularisation**, comme :

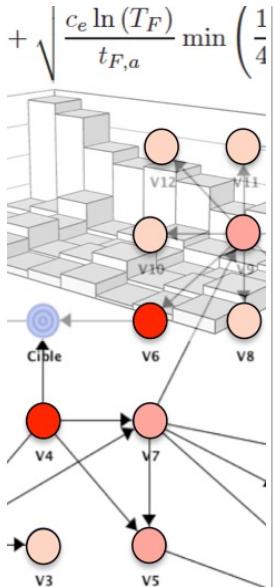
- Augmenter artificiellement le jeu de données; dans le cas de l'images on applique des opérateurs de transformation (rotation, translation, zoom, ...)



- Dropout : désactiver une certaine proportion de neurones pendant la phase d'apprentissage



Deep Learning - Transfer Learning



```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.  
        
```

$$= \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

Stéphane BONNEVAY – Polytech Lyon

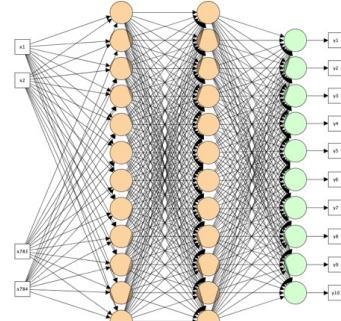
Classification d'images - Réseaux multicouches

Dropout avec Keras

```
model3 = Sequential()  
model3.add(Dense(50, input_dim=784, activation='relu'))  
model3.add(Dense(50, activation='relu'))  
model3.add(Dense(10, activation='softmax'))  
model3.summary()
```

Model: "sequential_3"

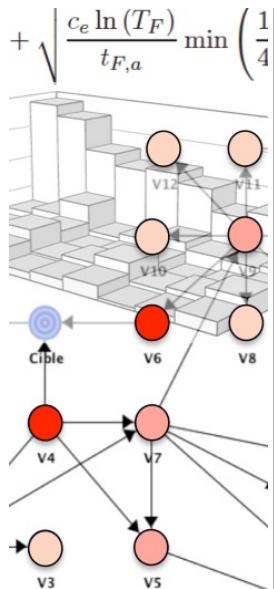
Layer (type)	Output Shape	Param #
dense_6 (Dense)	(None, 50)	39250
dense_7 (Dense)	(None, 50)	2550
dense_8 (Dense)	(None, 10)	510
Total params: 42,310		
Trainable params: 42,310		
Non-trainable params: 0		



```
model4 = Sequential()  
model4.add(Dense(50, input_dim=784, activation='relu'))  
model4.add(Dropout(0.5))  
model4.add(Dense(50, activation='relu'))  
model4.add(Dropout(0.5))  
model4.add(Dense(10, activation='softmax'))  
model4.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
dense_3 (Dense)	(None, 50)	39250
dropout_2 (Dropout)	(None, 50)	0
dense_4 (Dense)	(None, 50)	2550
dropout_3 (Dropout)	(None, 50)	0
dense_5 (Dense)	(None, 10)	510
Total params: 42,310		
Trainable params: 42,310		
Non-trainable params: 0		



```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.
```

$$+ \sqrt{\frac{c_e \ln(T_F)}{t_{F,a}}} \min \left(\frac{1}{4}$$

1. Introduction

2. Classification d'images avec

- un réseau multi-couches
- un réseau de neurones convolutifs
- un réseau pré-entraîné (transfer learning)

3. Détection d'objets dans une image

Deep Learning - Transfer Learning

Classification d'images - CNN

$$+ \sqrt{\frac{c_e \ln(T_F)}{t_{F,a}}} \min\left(\frac{1}{4}, \dots\right)$$

```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.  
        ...  
        :  $\prod_{j=1}^{q_i} \left( \frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$ 
```

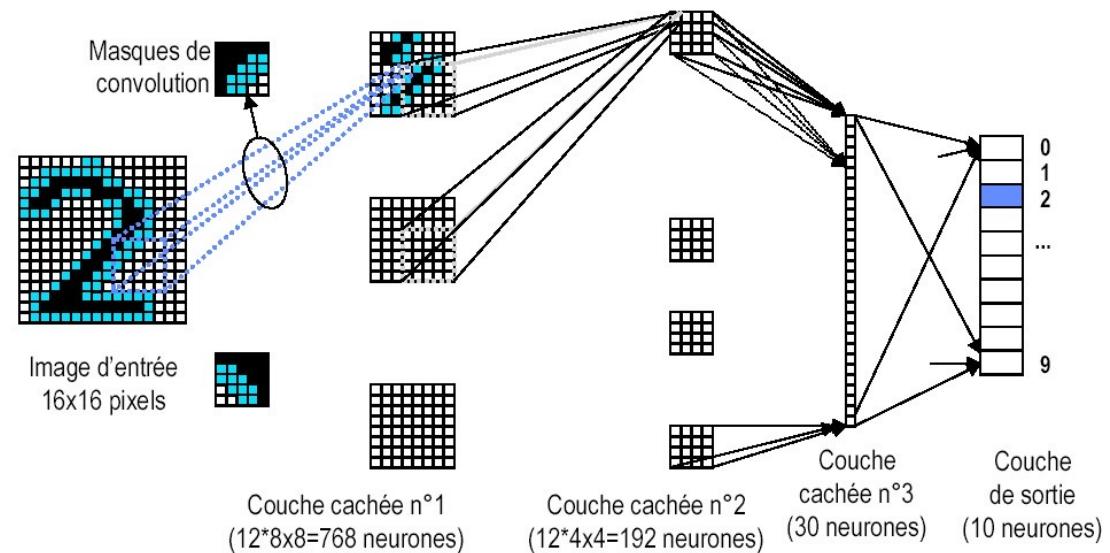


Matrice de taille 28x28 en entrée

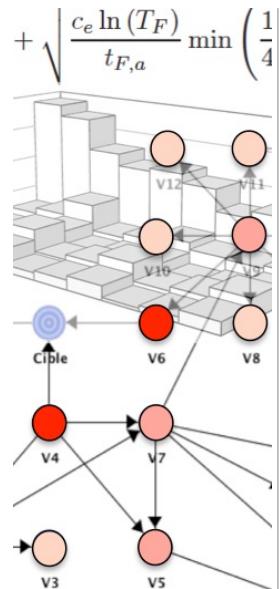
Conserver la structure de l'image

Extraire des propriétés liées à la topologie ← Masques de convolution

« LeNet » LeCun 1989



Deep Learning - Transfer Learning



```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.  
    }
```

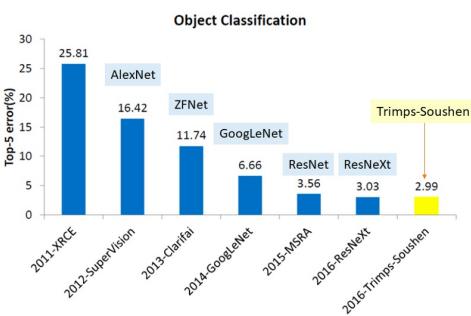
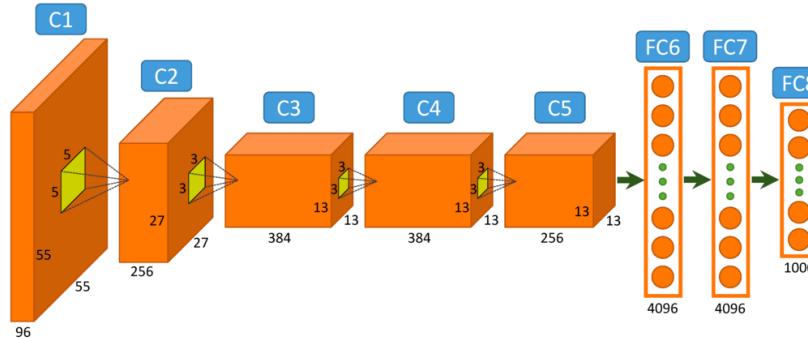
$$= \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

Classification d'images - CNN

Même si les réseaux de neurones convolutifs sont nés à la fin des années 80, en particulier avec « LeNet » de Yann Lecun, il faut attendre 2012 pour que la notion de Deep learning prenne tout son sens.

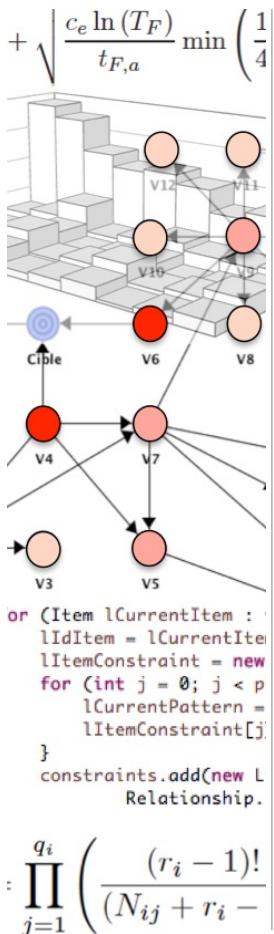
Lors de la compétition de reconnaissance d'images **IMAGENET** de 2012, des chercheurs ont explosé le record de reconnaissance grâce à une technique efficace de rétropropagation du gradient mais aussi l'usage de grosses puissances de calcul.

<https://www.image-net.org/challenges/LSVRC/>



Depuis, des équipes ont développés des réseaux avec plusieurs millions de neurones et donc plusieurs milliards de paramètres ... et beaucoup de données !

Deep Learning - Transfer Learning

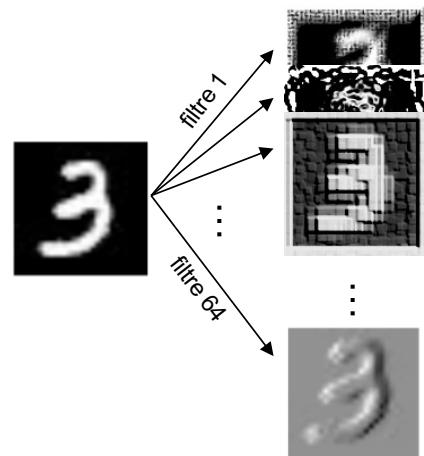
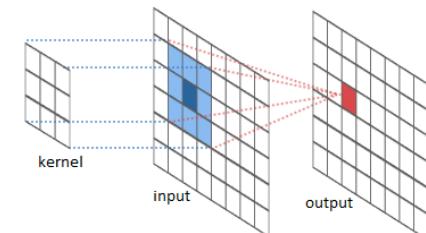


Classification d'images - CNN

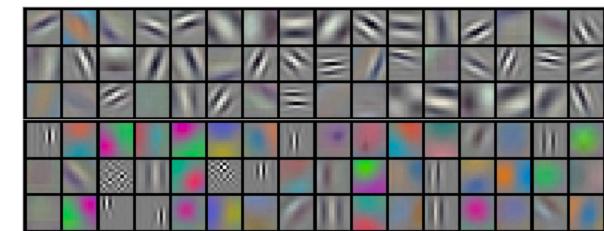
Masques de convolution

Diagram illustrating a convolution operation. An input image (28x28) is multiplied by a filter (3x3) to produce an output image (26x26 or 28x28 with padding).

image 28x28 * filtre 3x3 = image résultat 26x26 ou 28x28 avec padding

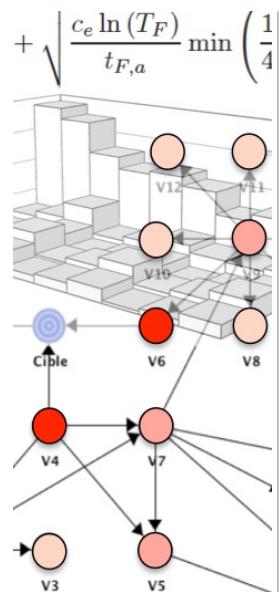


La première couche d'un réseau convolutif consiste à apprendre automatiquement, lors de la phase d'apprentissage, les valeurs des filtres.



Deep Learning - Transfer Learning

Classification d'images - CNN



```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.
```

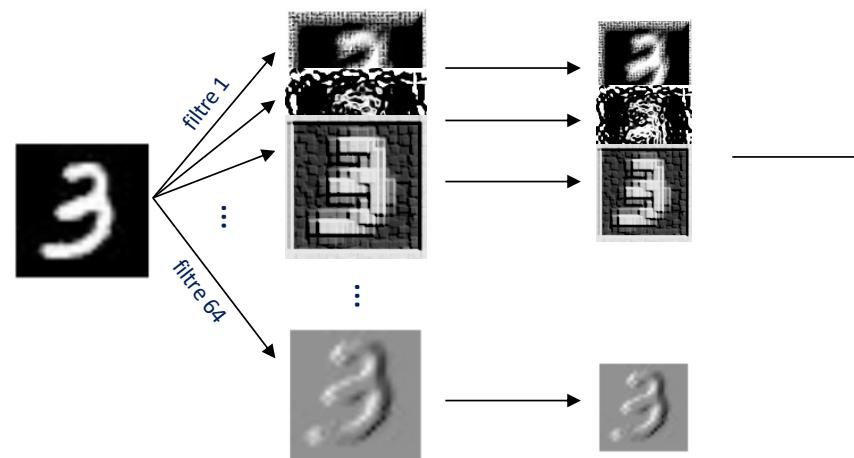
$$= \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

Max pooling : diminuer la taille du réseau

12	20	30	0
8	12	2	0
34	70	37	4
112	100	25	12

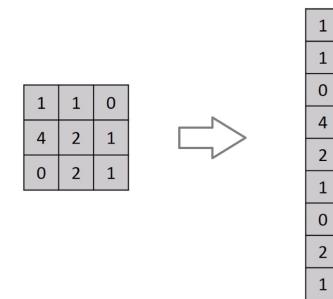
2 × 2 Max-Pool →

20	30
112	37

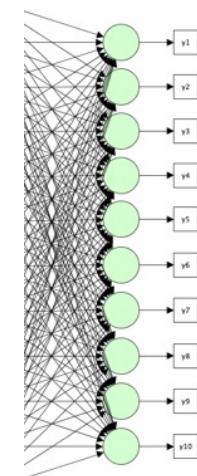


Stéphane BONNEVAY – Polytech Lyon

Flatten :



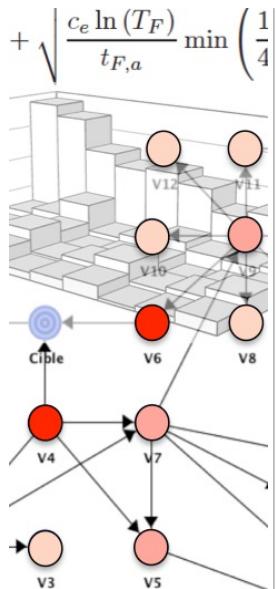
23
12
18
14
5
32
24
87
54
...
65
43
18
1
90
1
12
3
8
34



Deep Learning - Transfer Learning

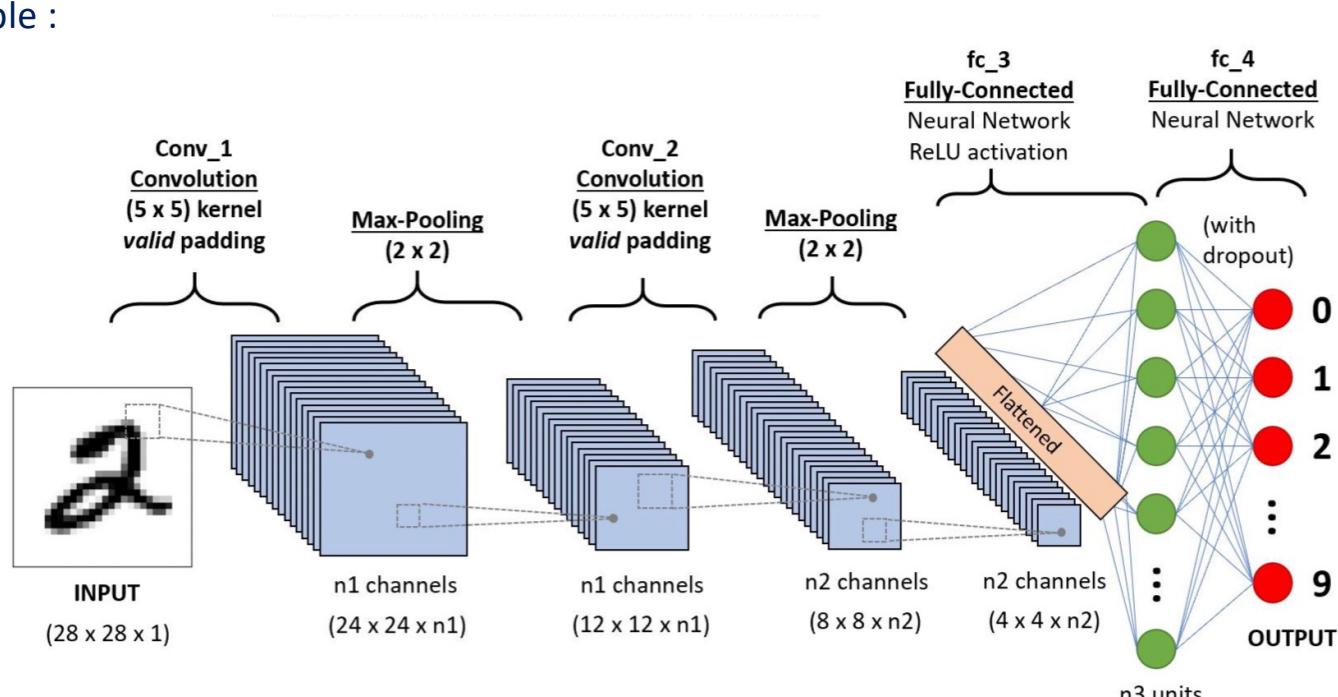
Classification d'images - CNN

Exemple :



```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.
```

$$= \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$



Source : <https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53>

Deep Learning - Transfer Learning

$$+ \sqrt{\frac{c_e \ln(T_F)}{t_{F,a}}} \min\left(\frac{1}{4}, \dots\right)$$

```

or (Item lCurrentItem :
    lIdItem = lCurrentItem
    lItemConstraint = new
    for (int j = 0; j < p
        lCurrentPattern =
        lItemConstraint[j];
    }
    constraints.add(new L
        Relationship.
    )
    :  $\prod_{j=1}^{q_i} \left( \frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$ 

```

Classification d'images - CNN



Matrice de taille 28x28 en entrée

$54\ 000 \times 28 \times 28 \times 1$
 $6\ 000 \times 28 \times 28 \times 1$
 $10\ 000 \times 28 \times 28 \times 1$

1 channel car niveau de gris

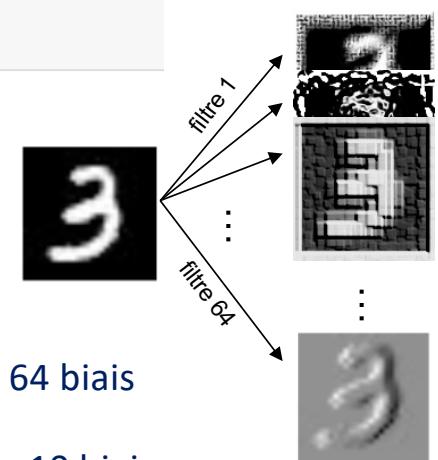
```

model4 = Sequential()
model4.add(Conv2D(filters=64, kernel_size=2, padding='same', activation='relu', input_shape=(28,28, 1)))
model4.add(MaxPooling2D(pool_size=2))
model4.add(Flatten())
model4.add(Dense(10, activation='softmax'))
model4.summary()

Model: "sequential_4"
Layer (type)          Output Shape       Param #
conv2d (Conv2D)        (None, 28, 28, 64)   320
max_pooling2d (MaxPooling2D) (None, 14, 14, 64)   0
flatten (Flatten)      (None, 12544)        0
dense_9 (Dense)        (None, 10)           125450
=====
Total params: 125,770
Trainable params: 125,770
Non-trainable params: 0

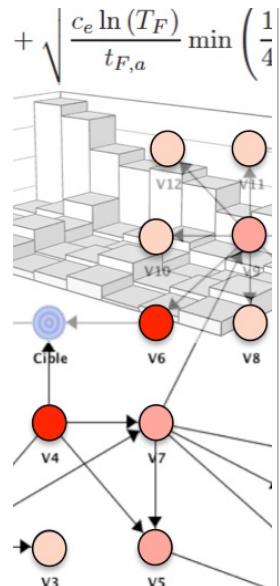
```

$64 \times (2 \times 2) + 64$ biais
 $12544 \times 10 + 10$ biais



Deep Learning - Transfer Learning

Classification d'images - CNN



```

or (Item lCurrentItem :
    lIdItem = lCurrentItem
    lItemConstraint = new
    for (int j = 0; j < p
        lCurrentPattern =
        lItemConstraint[j]
    }
    constraints.add(new L
        Relationship.

```

$$= \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

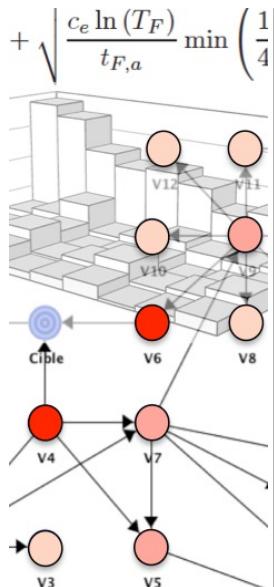
```

Train on 54000 samples, validate on 6000 samples
Epoch 1/10
54000/54000 [=====] - 19s 357us/sample - loss: 0.4259 - accuracy: 0.8526 - val_loss: 0.3415
- val_accuracy: 0.8795
Epoch 2/10
54000/54000 [=====] - 20s 370us/sample - loss: 0.3069 - accuracy: 0.8911 - val_loss: 0.3116
- val_accuracy: 0.8907
Epoch 3/10
54000/54000 [=====] - 19s 345us/sample - loss: 0.2766 - accuracy: 0.9022 - val_loss: 0.2978
- val_accuracy: 0.8932
Epoch 4/10
54000/54000 [=====] - 20s 376us/sample - loss: 0.2557 - accuracy: 0.9082 - val_loss: 0.2981
- val_accuracy: 0.8922
Epoch 5/10
54000/54000 [=====] - 20s 374us/sample - loss: 0.2408 - accuracy: 0.9139 - val_loss: 0.2831
- val_accuracy: 0.8985
Epoch 6/10
54000/54000 [=====] - 20s 378us/sample - loss: 0.2251 - accuracy: 0.9194 - val_loss: 0.2674
- val_accuracy: 0.9040
Epoch 7/10
54000/54000 [=====] - 21s 392us/sample - loss: 0.2119 - accuracy: 0.9240 - val_loss: 0.2743
- val_accuracy: 0.9013
Epoch 8/10
54000/54000 [=====] - 21s 387us/sample - loss: 0.2020 - accuracy: 0.9276 - val_loss: 0.2663
- val_accuracy: 0.9075
Epoch 9/10
54000/54000 [=====] - 21s 394us/sample - loss: 0.1920 - accuracy: 0.9322 - val_loss: 0.2947
- val_accuracy: 0.8958
Epoch 10/10
54000/54000 [=====] - 21s 393us/sample - loss: 0.1838 - accuracy: 0.9341 - val_loss: 0.2746
- val_accuracy: 0.9035

```

90,4% de précision sur les données de validation

90,2% de précision sur les données de test

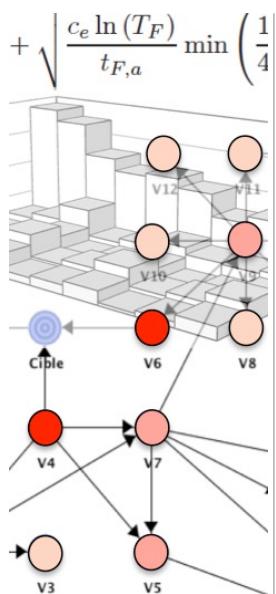


```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.
```

$$= \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

Deep Learning - Transfer Learning

Classification d'images - Transfer Learning

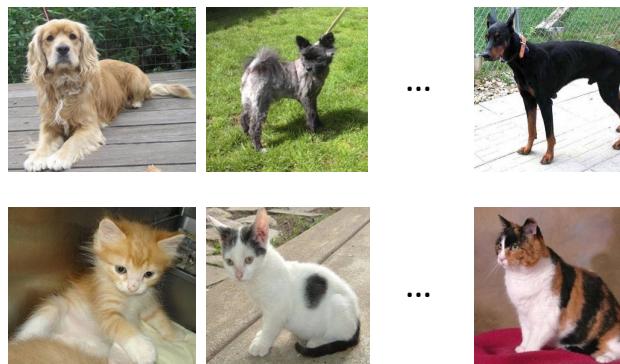


```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.  
        ...)
```

$$\prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

Problème

Reconnaissance et catégorisation d'images en deux catégories : chat ou chien



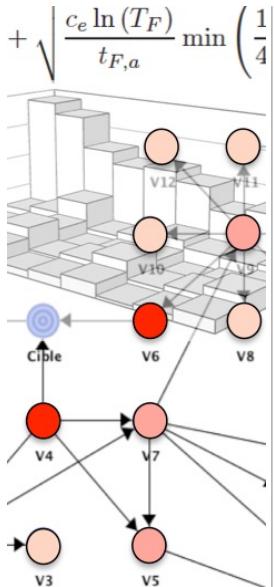
<https://www.kaggle.com/c/dogs-vs-cats/data>

Images en couleur
de taille 150x150

Contraintes

- On ne dispose que de 3000 images (1500 de chaque catégorie) pour l'apprentissage et 1000 images (500 de chaque catégorie) pour la validation et le test
- On ne dispose pas d'une grosse puissance de calculs

Deep Learning - Transfer Learning



```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.
```

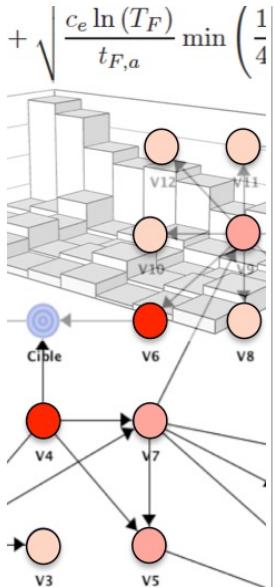
$$:= \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

Classification d'images - Transfer Learning

Apprentissage du CNN à partir des 3000 images dédiées à l'apprentissage et des 1000 images pour la validation (les images sont normalisées entre 0 et 1).

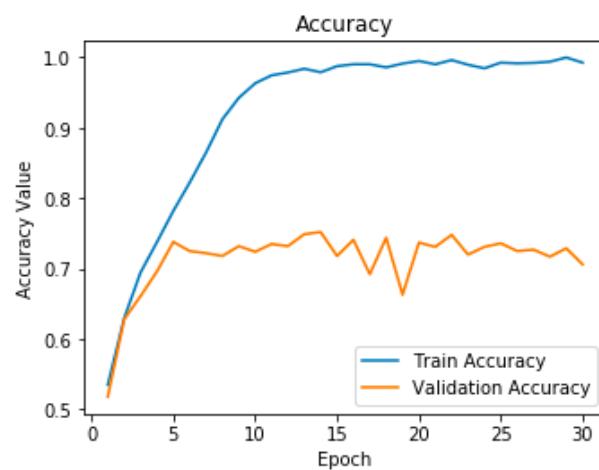
Layer (type)	Output Shape	Param #
<hr/>		
conv2d_1 (Conv2D)	(None, 148, 148, 16)	448
max_pooling2d_1 (MaxPooling2)	(None, 74, 74, 16)	0
conv2d_2 (Conv2D)	(None, 72, 72, 64)	9280
max_pooling2d_2 (MaxPooling2)	(None, 36, 36, 64)	0
conv2d_3 (Conv2D)	(None, 34, 34, 128)	73856
max_pooling2d_3 (MaxPooling2)	(None, 17, 17, 128)	0
flatten_1 (Flatten)	(None, 36992)	0
dense_1 (Dense)	(None, 512)	18940416
dense_2 (Dense)	(None, 1)	513
<hr/>		
Total params: 19,024,513		
Trainable params: 19,024,513		
Non-trainable params: 0		

Deep Learning - Transfer Learning

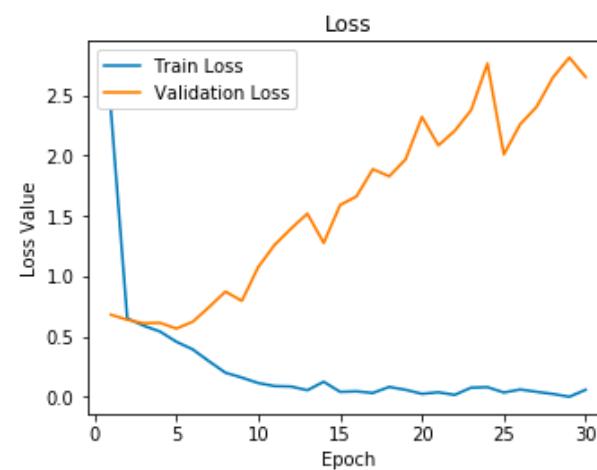


```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.
```

$$: \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

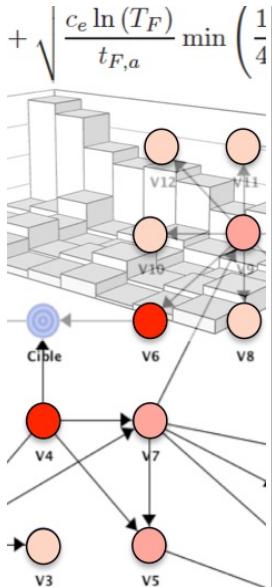


Classification d'images - Transfer Learning



On observe un **surapprentissage** après 3 epochs
Néanmoins, la précision moyenne sur les données de validation est de l'ordre de **71%**

Deep Learning - Transfer Learning



```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.
```

$$: \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

Stéphane BONNEVAY – Polytech Lyon

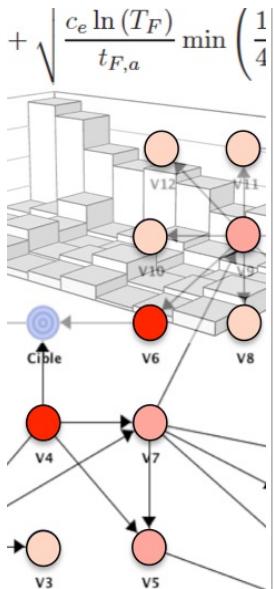
Ajout d'une couche supplémentaire mais surtout ajout d'un processus de **dropout** à 0.3 après chaque couche cachée dense

Apprentissage du CNN à partir des 3000 images dédiées à l'apprentissage et des 1000 images pour la validation (les images sont normalisées entre 0 et 1).

Classification d'images - Transfer Learning

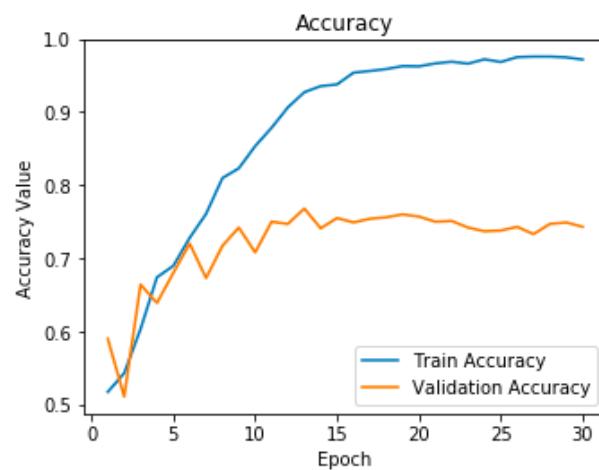
Layer (type)	Output Shape	Param #
<hr/>		
conv2d_1 (Conv2D)	(None, 148, 148, 16)	448
max_pooling2d_1 (MaxPooling2)	(None, 74, 74, 16)	0
conv2d_2 (Conv2D)	(None, 72, 72, 64)	9280
max_pooling2d_2 (MaxPooling2)	(None, 36, 36, 64)	0
conv2d_3 (Conv2D)	(None, 34, 34, 128)	73856
max_pooling2d_3 (MaxPooling2)	(None, 17, 17, 128)	0
conv2d_4 (Conv2D)	(None, 15, 15, 128)	147584
max_pooling2d_4 (MaxPooling)	(None, 7, 7, 128)	0
flatten_1 (Flatten)	(None, 6272)	0
dense_2 (Dense)	(None, 512)	3211776
dropout_1 (Dropout)	(None, 512)	0
dense_3 (Dense)	(None, 512)	262656
dropout_2 (Dropout)	(None, 512)	0
dense_4 (Dense)	(None, 1)	513
<hr/>		
Total params: 3,706,113		
Trainable params: 3,706,113		
Non-trainable params: 0		

Deep Learning - Transfer Learning

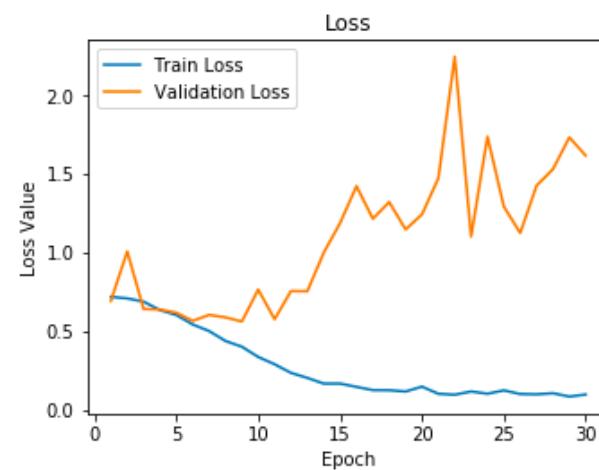


```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.
```

Stéphane BONNEVAY – Polytech Lyon



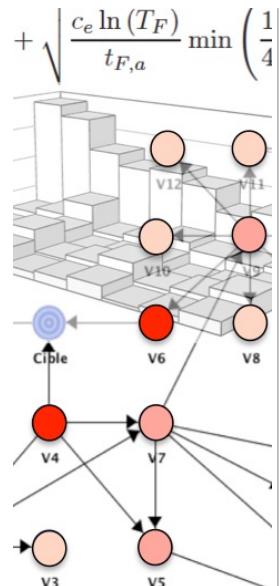
Classification d'images - Transfer Learning



Il y a toujours un **surapprentissage** mais après plus d'epochs
Néanmoins, la précision moyenne sur les données de validation est un peu plus élevée de
l'ordre de **74%**

Deep Learning - Transfer Learning

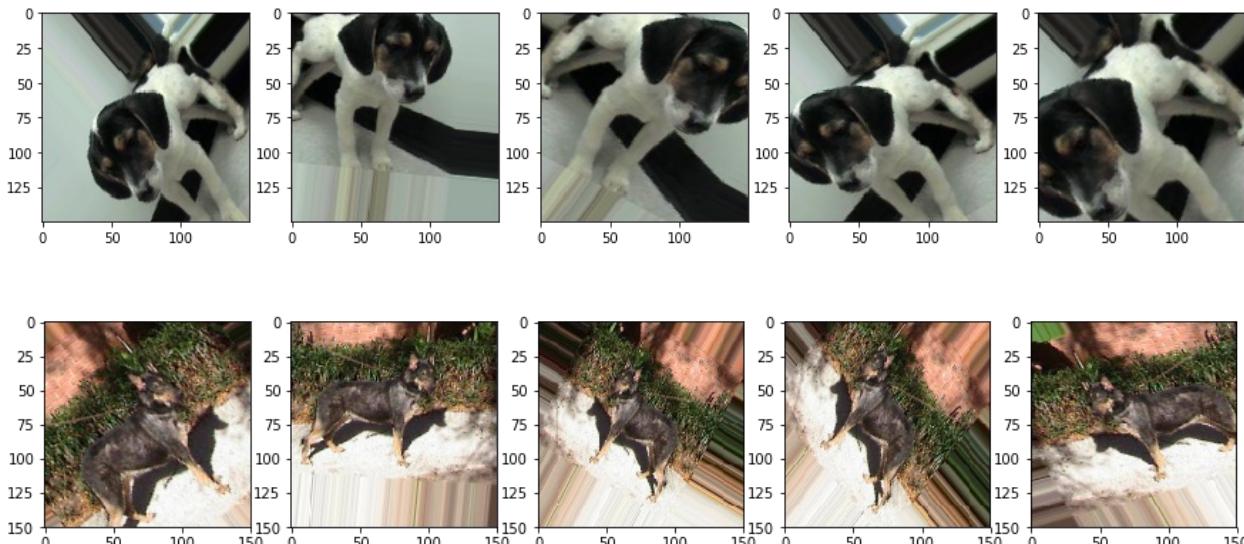
Classification d'images - Transfer Learning



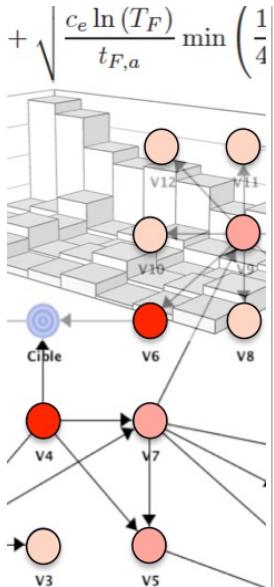
```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.  
        ...)
```

$$= \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

L'idée est de créer de nouvelles images à partir des 3000 images d'apprentissage en appliquant des opérateurs de rotation, de translation, de zoom, ... afin d'augmenter le nombre d'images tout en produisant des Nouvelles.

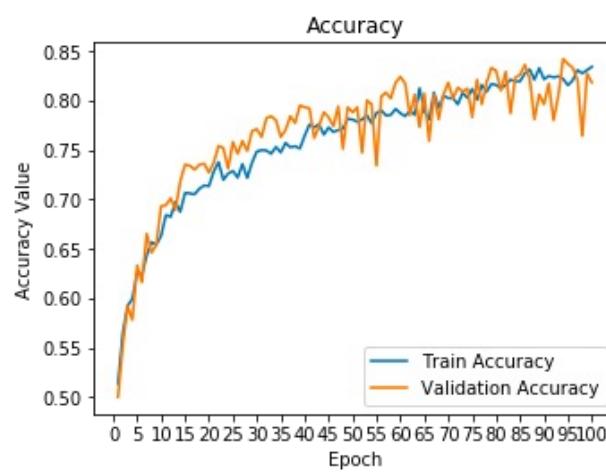


Deep Learning - Transfer Learning

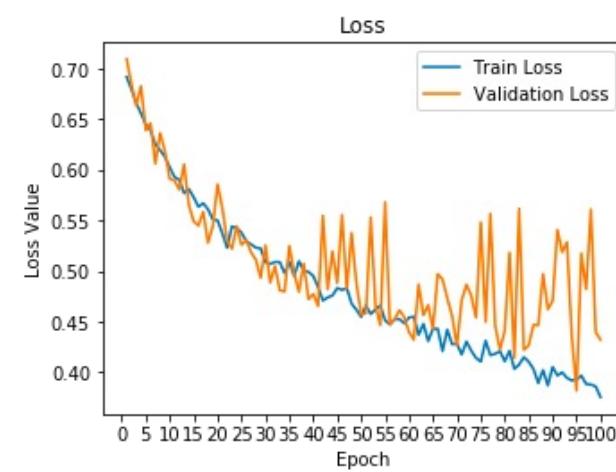


```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.
```

Stéphane BONNEVAY – Polytech Lyon



Classification d'images - Transfer Learning

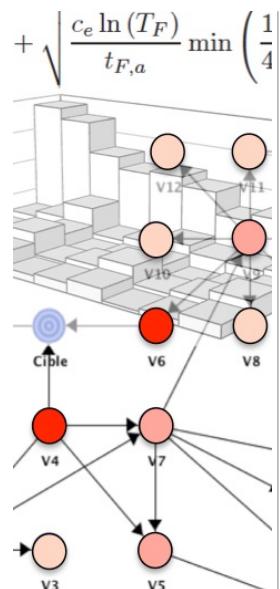


Il n'y a plus de surapprentissage !

La précision moyenne sur les données de validation est encore un peu élevée de l'ordre de 80%

Deep Learning - Transfer Learning

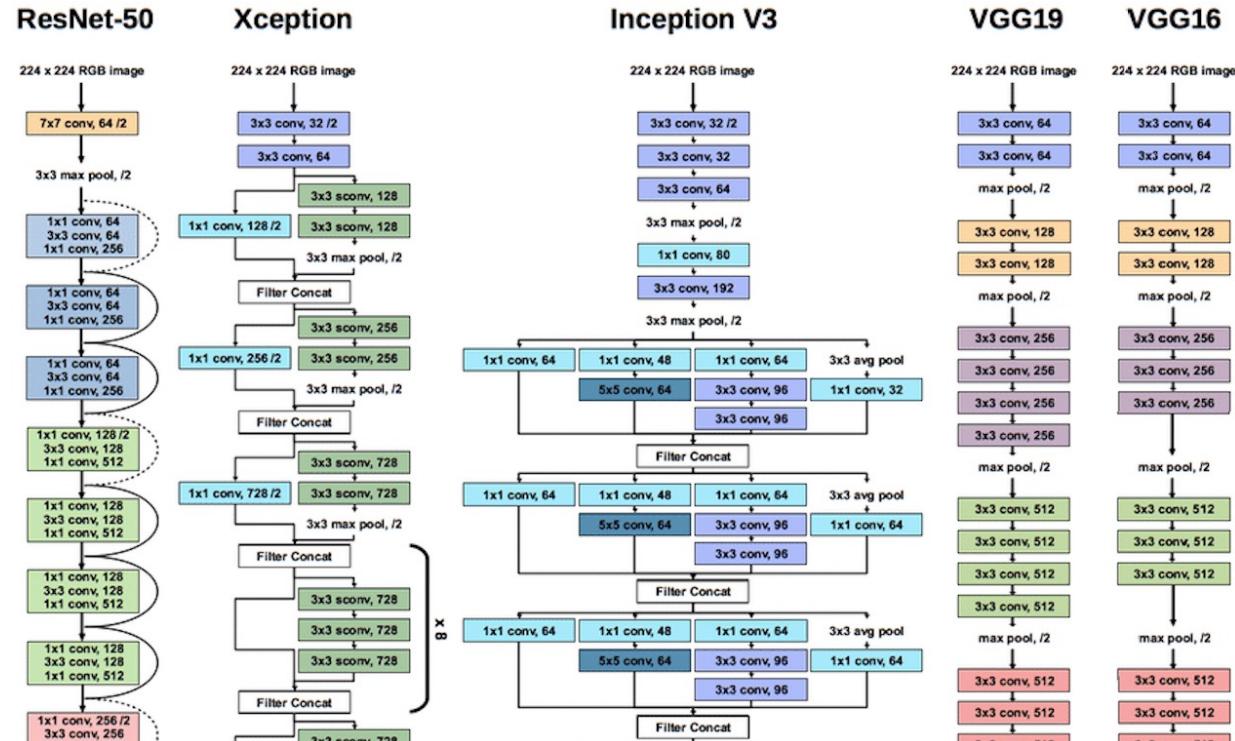
Classification d'images - Transfer Learning



```

or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.
    
```

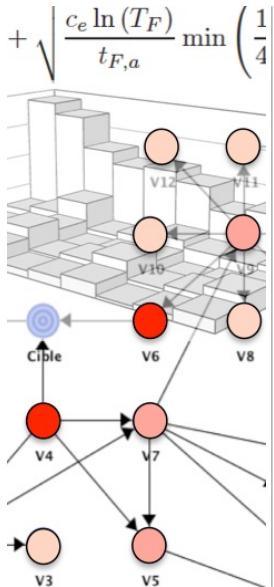
Stéphane BONNEVAY – Polytech Lyon



https://www.researchgate.net/publication/330478807_Deep_Feature-Based_Classifiers_for_Fruit_Fly_Identification_Diptera_Tephritidae/download

$$\prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

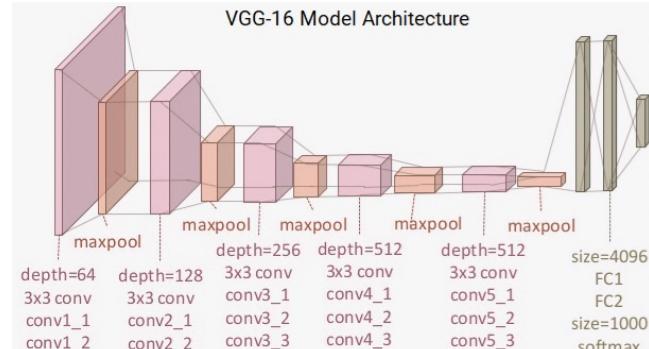
Deep Learning - Transfer Learning



```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.  
    )
```

$$: \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

Stéphane BONNEVAY – Polytech Lyon



Published as a conference paper at ICLR 2015

<https://arxiv.org/abs/1409.1556>

VERY DEEP CONVOLUTIONAL NETWORKS FOR LARGE-SCALE IMAGE RECOGNITION

Karen Simonyan* & Andrew Zisserman*
Visual Geometry Group, Department of Engineering Science, University of Oxford
{karen,az}@robots.ox.ac.uk

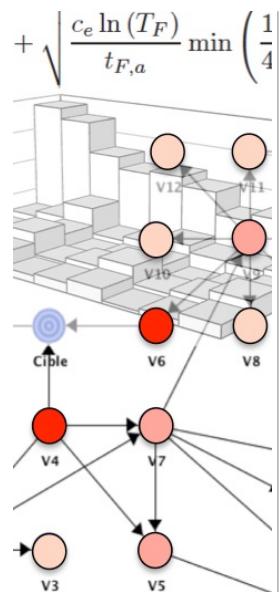
Apprentissage à partir de 1.3M d'images
50K images pour la validation
100K images pour les tests

Classification d'images - Transfer Learning

Layer (type)	Output Shape	Param #
<hr/>		
input_1 (InputLayer)	(None, 150, 150, 3)	0
block1_conv1 (Conv2D)	(None, 150, 150, 64)	1792
block1_conv2 (Conv2D)	(None, 150, 150, 64)	36928
block1_pool (MaxPooling2D)	(None, 75, 75, 64)	0
block2_conv1 (Conv2D)	(None, 75, 75, 128)	73856
block2_conv2 (Conv2D)	(None, 75, 75, 128)	147584
block2_pool (MaxPooling2D)	(None, 37, 37, 128)	0
block3_conv1 (Conv2D)	(None, 37, 37, 256)	295168
block3_conv2 (Conv2D)	(None, 37, 37, 256)	590080
block3_conv3 (Conv2D)	(None, 37, 37, 256)	590080
block3_pool (MaxPooling2D)	(None, 18, 18, 256)	0
block4_conv1 (Conv2D)	(None, 18, 18, 512)	1180160
block4_conv2 (Conv2D)	(None, 18, 18, 512)	2359808
block4_conv3 (Conv2D)	(None, 18, 18, 512)	2359808
block4_pool (MaxPooling2D)	(None, 9, 9, 512)	0
block5_conv1 (Conv2D)	(None, 9, 9, 512)	2359808
block5_conv2 (Conv2D)	(None, 9, 9, 512)	2359808
block5_conv3 (Conv2D)	(None, 9, 9, 512)	2359808
block5_pool (MaxPooling2D)	(None, 4, 4, 512)	0
flatten_6 (Flatten)	(None, 8192)	0
<hr/>		
Total params:	14,714,688	
Trainable params:	0	
Non-trainable params:	14,714,688	

Deep Learning - Transfer Learning

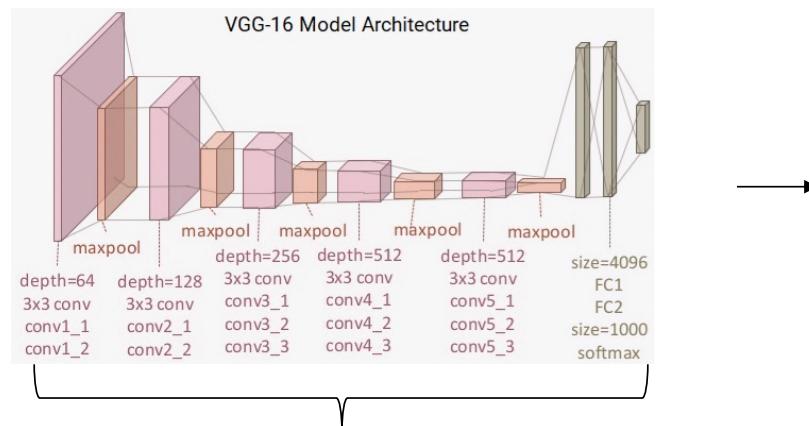
Classification d'images - Transfer Learning



```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.  
    )
```

$$= \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

Modèle VGG-16 utilisé comme extracteur de features



Extraction de caractéristiques
modèle pré-entraîné inchangé

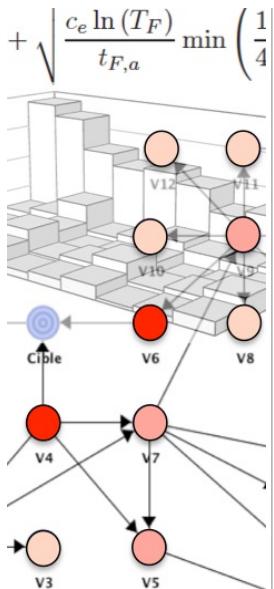
Layer (type)	Output Shape	Param #
<hr/>		
dense_14 (Dense)	(None, 512)	4194816
<hr/>		
dropout_7 (Dropout)	(None, 512)	0
<hr/>		
dense_15 (Dense)	(None, 512)	262656
<hr/>		
dropout_8 (Dropout)	(None, 512)	0
<hr/>		
dense_16 (Dense)	(None, 1)	513
<hr/>		
Total params: 4,457,985		
Trainable params: 4,457,985		
Non-trainable params: 0		

Classification
modèle à entraîner

Intérêt :

- Profiter d'un réseau plus efficace, en terme d'extraction de caractéristiques, qu'un réseau qu'on aurait construit nous même avec nos propres données
- Ne réaliser l'apprentissage que sur quelques couches (compatible avec notre puissance de calculs et notre quantité de données)

Deep Learning - Transfer Learning

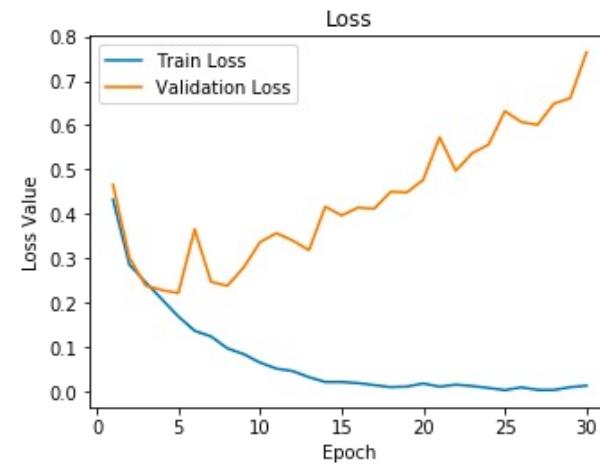
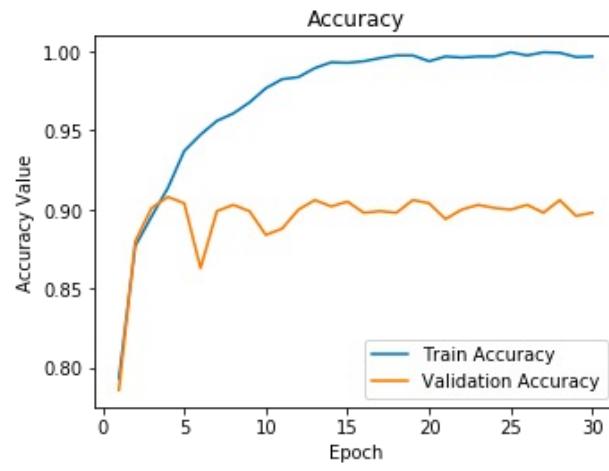


```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.
```

$$= \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

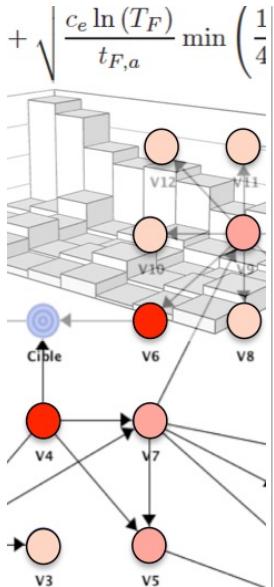
Classification d'images - Transfer Learning

Modèle VGG-16 utilisé comme extracteur de features



Il y a de nouveau du surapprentissage
Par contre, la précision moyenne sur les données de validation est de l'ordre de 88%

Deep Learning - Transfer Learning

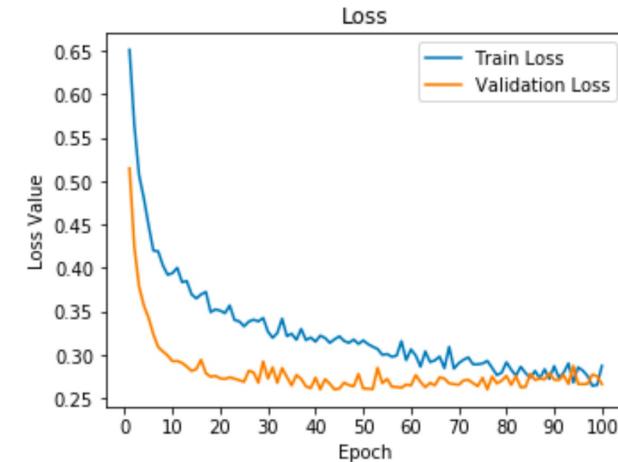
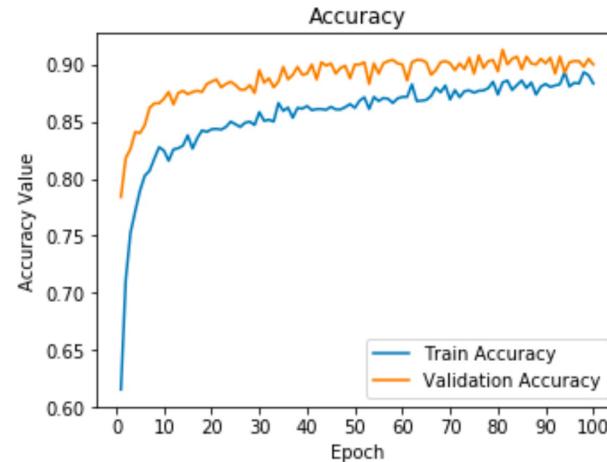


```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.
```

$$= \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

Classification d'images - Transfer Learning

Modèle VGG-16 utilisé comme extracteur de features et + d'images

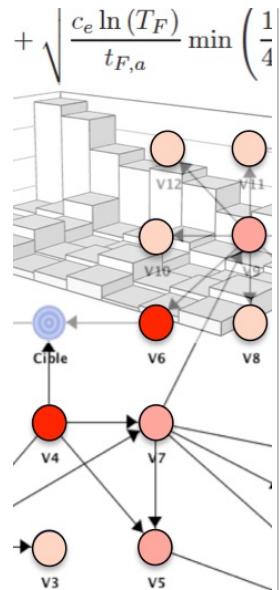


Il n'y a plus de surapprentissage

La précision moyenne sur les données de validation est de l'ordre de 90%

Deep Learning - Transfer Learning

Classification d'images - Transfer Learning



```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.  
    }
```

$$= \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

Stéphane BONNEVAY – Polytech Lyon

Modèle VGG-16 utilisé comme extracteur de features et + d'images et Fine-Tuning

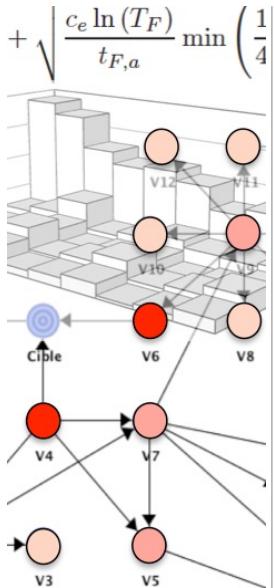
Layer Type	Layer Name	Layer Trainable
0	input_1	False
1	block1_conv1	False
2	block1_conv2	False
3	block1_pool	False
4	block2_conv1	False
5	block2_conv2	False
6	block2_pool	False
7	block3_conv1	False
8	block3_conv2	False
9	block3_conv3	False
10	block3_pool	False
11	block4_conv1	True
12	block4_conv2	True
13	block4_conv3	True
14	block4_pool	True
15	block5_conv1	True
16	block5_conv2	True
17	block5_conv3	True
18	block5_pool	True
19	flatten_6	True

Layer (type)	Output Shape	Param #
model_1 (Model)	(None, 8192)	14714688
dense_23 (Dense)	(None, 512)	4194816
dropout_13 (Dropout)	(None, 512)	0
dense_24 (Dense)	(None, 512)	262656
dropout_14 (Dropout)	(None, 512)	0
dense_25 (Dense)	(None, 1)	513

Total params: 19,172,673
Trainable params: 17,437,185
Non-trainable params: 1,735,488

Classification

Deep Learning - Transfer Learning

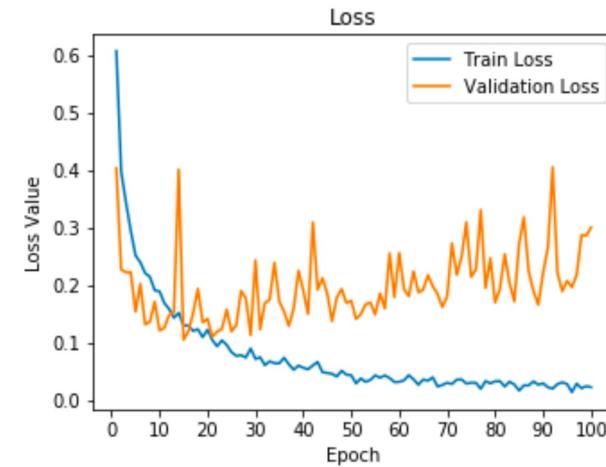
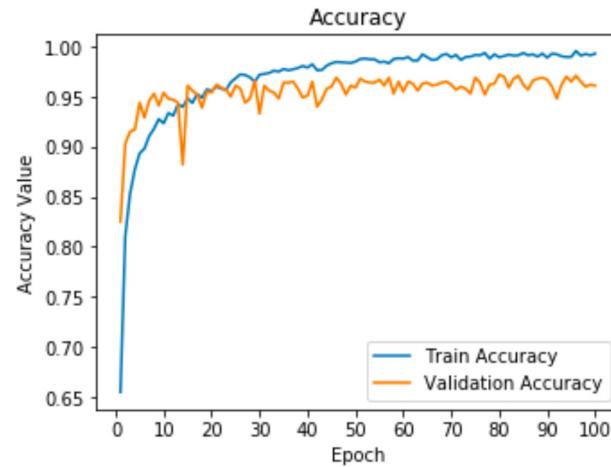


```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.
```

$$: \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

Classification d'images - Transfer Learning

Modèle VGG-16 utilisé comme extracteur de features et + d'images et Fine-Tuning

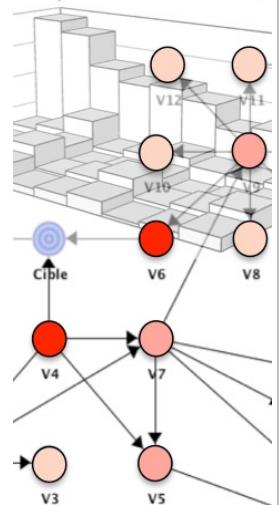


La précision moyenne sur les données de validation est maintenant de l'ordre de 96%

Deep Learning - Transfer Learning

Classification d'images - Transfer Learning

$$+ \sqrt{\frac{c_e \ln(T_F)}{t_{F,a}}} \min \left(\frac{1}{4}$$



```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.  
    )
```

$$= \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

Stéphane BONNEVAY – Polytech Lyon

Plusieurs modèles, dédiés à la classification d'images, ont été construits sur ImageNet et mis à disposition (Google Inception model, Microsoft ResNet model, Oxford VGG model, ...).

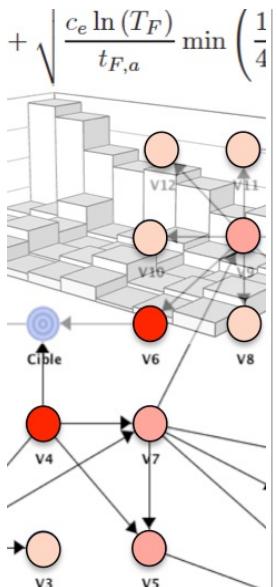
Model	Size (MB)	Top-1 Accuracy	Top-5 Accuracy	Parameters	Depth	Time (ms) per inference step (CPU)	Time (ms) per inference step (GPU)
Xception	88	79.0%	94.5%	22.9M	81	109.4	8.1
VGG16	528	71.3%	90.1%	138.4M	16	69.5	4.2
VGG19	549	71.3%	90.0%	143.7M	19	84.8	4.4
ResNet50	98	74.9%	92.1%	25.6M	107	58.2	4.6
ResNet50V2	98	76.0%	93.0%	25.6M	103	45.6	4.4
ResNet101	171	76.4%	92.8%	44.7M	209	89.6	5.2
ResNet101V2	171	77.2%	93.8%	44.7M	205	72.7	5.4
ResNet152	232	76.6%	93.1%	60.4M	311	127.4	6.5
ResNet152V2	232	78.0%	94.2%	60.4M	307	107.5	6.6
InceptionV3	92	77.9%	93.7%	23.9M	189	42	
InceptionResNetV2	215	80.3%	95.3%	55.9M	449	130	
MobileNet	16	70.4%	89.5%	4.3M	55	22	
MobileNetV2	14	71.3%	90.1%	2.5M	105	22	



```
keras.applications.inception_v3.InceptionV3(include_top=True,  
weights='imagenet', input_tensor=None,  
input_shape=None, pooling=None, classes=1000)
```

<https://keras.io/api/applications/>

Deep Learning - Transfer Learning



```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.  
    )
```

$$= \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

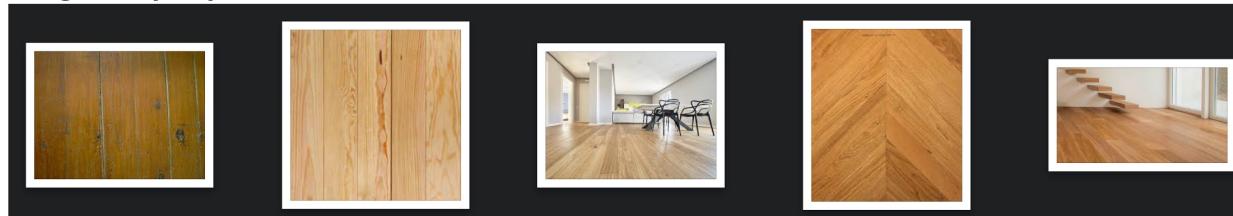
Classification d'images - Transfer Learning

Applications réelles :

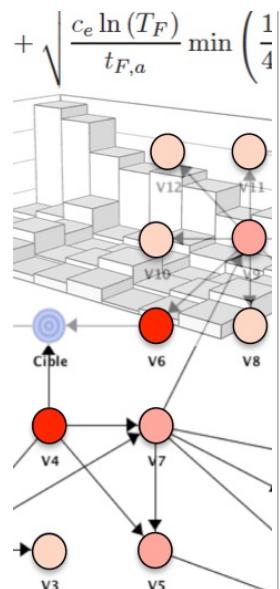
Images de carrelages :



Images de parquets :



Même avec seulement 500 images de chaque catégorie, on atteint 91% de précision moyenne.



```

or (Item lCurrentItem : 
    lIdItem = lCurrentItem
    lItemConstraint = new
    for (int j = 0; j < p
        lCurrentPattern =
        lItemConstraint[j];
    }
    constraints.add(new L
        Relationship.

```

$$= \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

Applications réelles :

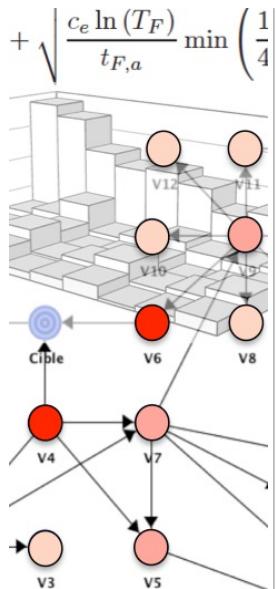
Images de “berline” :



Images de “pickup / 4x4” :



Avec 1000 images de chaque catégorie, on atteint seulement 74% de précision moyenne.

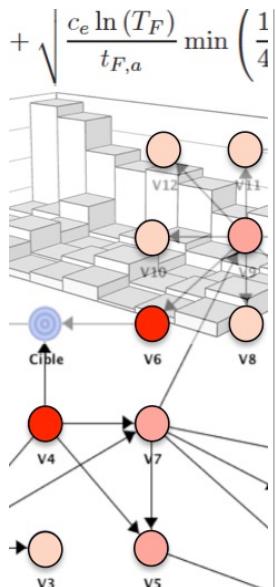


```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.  
        ...);  
}
```

$$= \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

Deep Learning - Transfer Learning

Détection d'objets



```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.  
        ...);
```

$$= \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

Classification

Chat

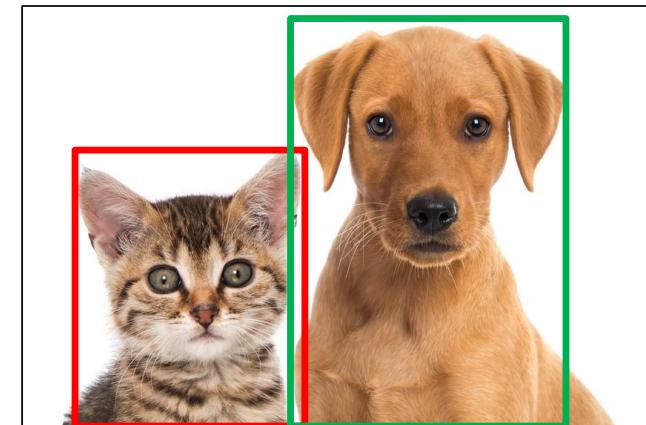


Chien

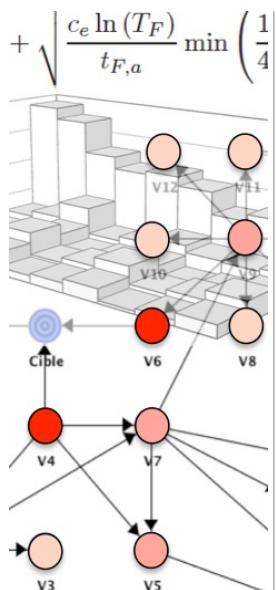


Détection

=
Localisation + Classification

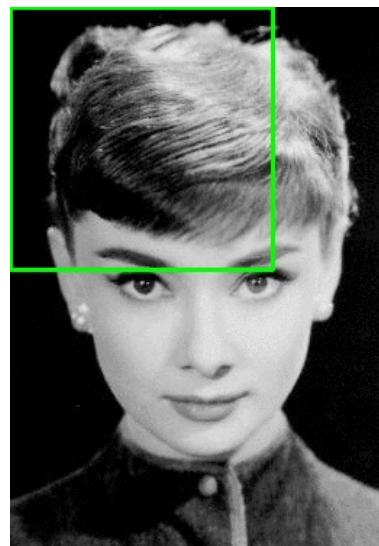


Pour une fenêtre glissante de taille donnée, la détection pourrait être une suite de classifications d'images obtenues en déplaçant la fenêtre dans l'image :

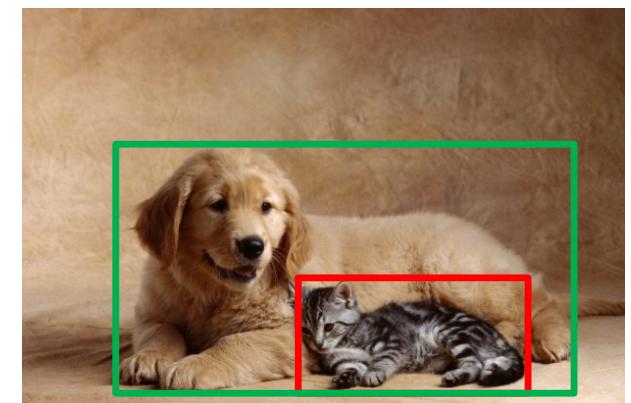


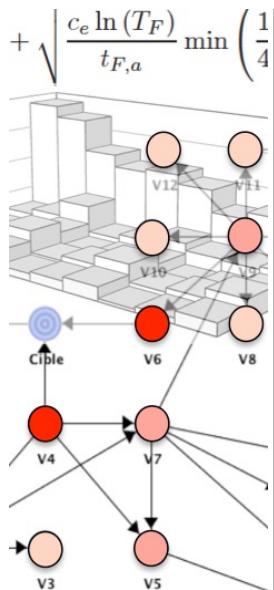
```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.  
        .
```

$$= \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$



Problème
les objets sont
de tailles variées





```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.  
    }
```

$$: \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

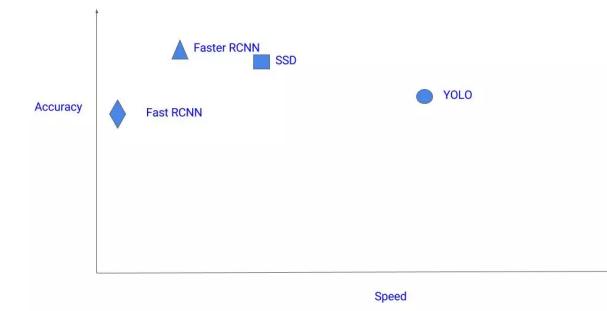
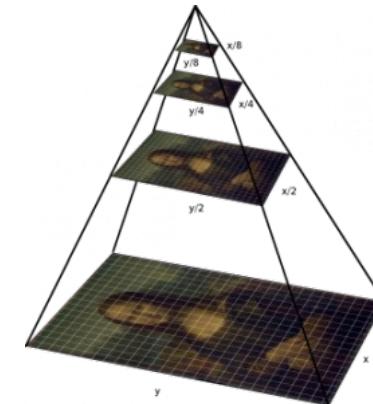
Idée :

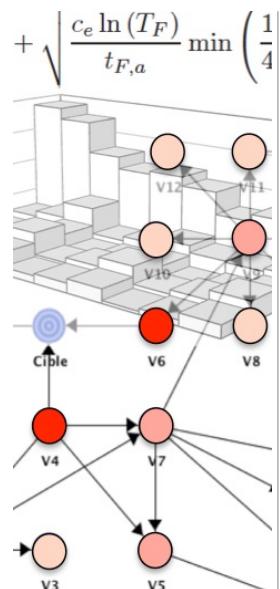
A partir de l'image originelle, construire une « pyramide » d'images à différentes échelles de plus en plus petites (généralement de 64 niveaux) et, avec une fenêtre glissante de taille fixe, parcourir ces différentes images pour détecter les objets.

Inconvénient : les objets n'ont pas tous la même forme

Il existe de nombreux algorithmes dont :

- Histogram of Oriented Gradients (HOG)
- Region-based Convolutional Neural Networks (R-CNN)
- Spatial Pyramid Pooling (SPP-net)
- Fast R-CNN
- Faster R-CNN
- You only Look Once (YOLO)
- Single Shot MultiBox Detector (SSD)
- RetinaNet





```

or (Item lCurrentItem :
    lIdItem = lCurrentItem
    lItemConstraint = new
    for (int j = 0; j < p
        lCurrentPattern =
        lItemConstraint[j];
    }
    constraints.add(new L
        Relationship.
    )
  
```

$$= \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

Stéphane BONNEVAY – Polytech Lyon

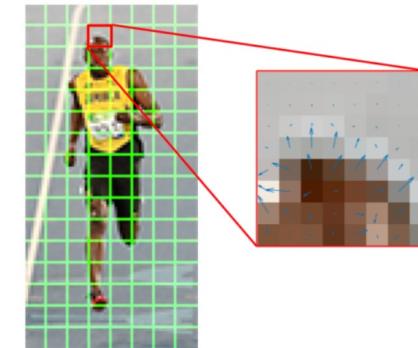
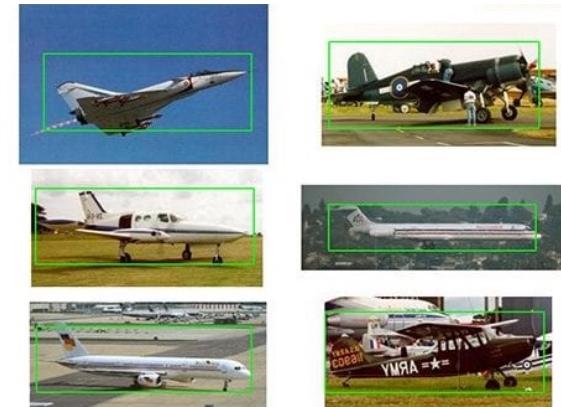
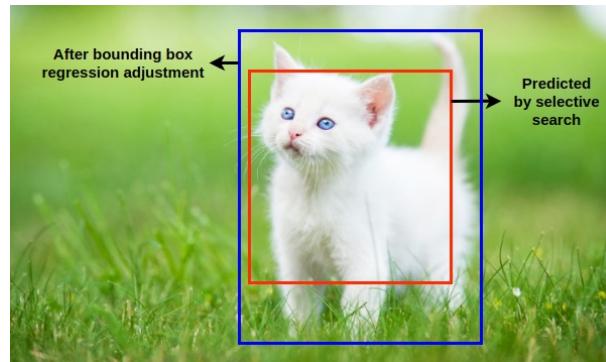
Histogram of Oriented Gradients (HOG)

Sur chaque fenêtre glissante :

- calcul de paramètres (Hog features)
- mis ensuite en entrée d'un SVM qui détermine la classe
- puis ajustement de la zone par :

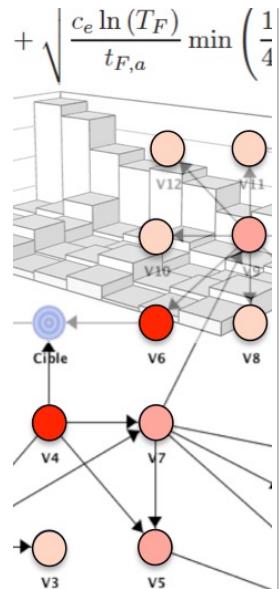
Navneet Dalal and Bill Triggs (2005)

Bounding box regression



Gradient Magnitude
2 3 4 4 3 4 2 2
5 11 17 13 7 9 3 4
11 21 23 27 22 17 4 6
23 99 165 135 85 32 26 2
91 155 133 136 144 152 57 28
98 196 76 38 26 60 170 51
165 60 60 27 77 85 43 136
71 13 34 23 108 27 48 110

Gradient Direction
80 36 5 10 0 64 90 73
37 9 9 179 78 27 169 166
87 136 173 39 102 163 152 176
76 13 1 168 159 22 125 143
120 70 14 150 145 144 145 143
58 86 119 98 100 101 133 113
30 65 157 75 78 165 145 124
11 170 91 4 110 17 133 110



```

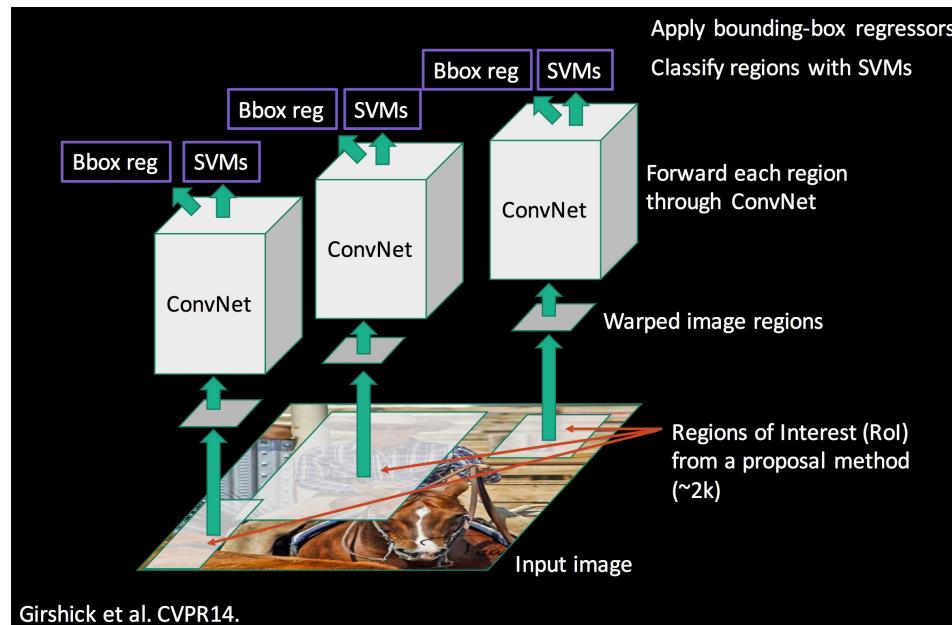
or (Item lCurrentItem :
    lIdItem = lCurrentItem
    lItemConstraint = new
    for (int j = 0; j < p
        lCurrentPattern =
        lItemConstraint[j];
    }
    constraints.add(new L
        Relationship.
    )
  
```

$$= \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

Stéphane BONNEVAY – Polytech Lyon

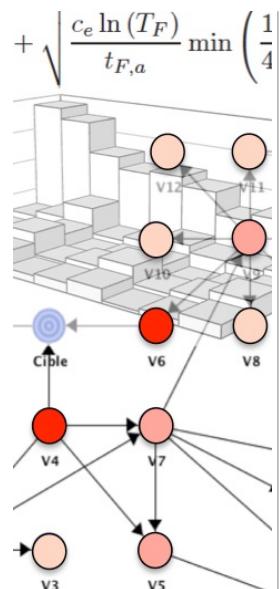
Region-based Convolutional Neural Networks (R-CNN)

Sélection de 2000 zones dans la pyramide (“Selective Search”) sur lesquelles les calculs de paramètres sont remplacés par un CNN, puis un SVM pour le classement :



Très lent !

Source : <https://arxiv.org/abs/1311.2524>

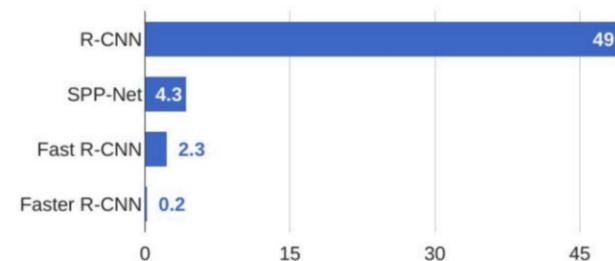
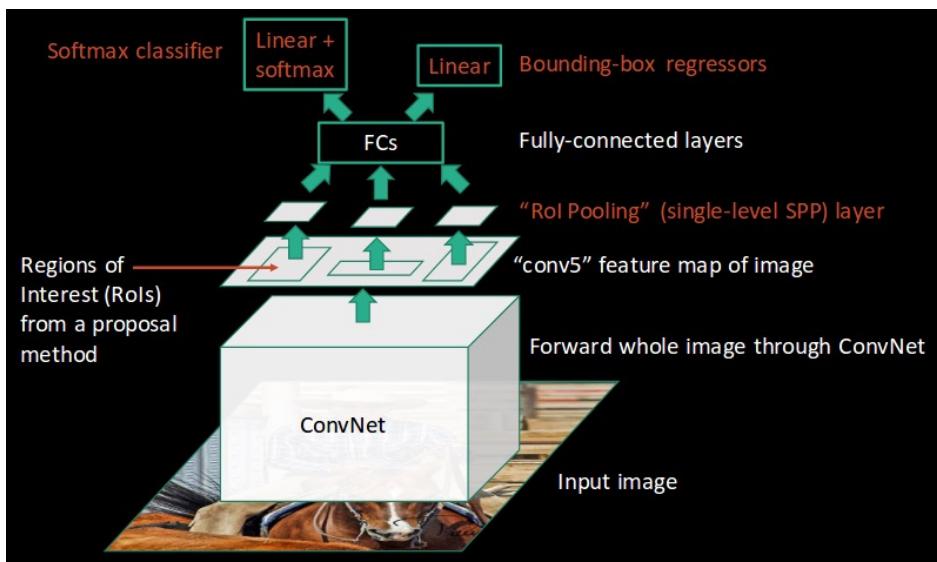


```

or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.
    
```

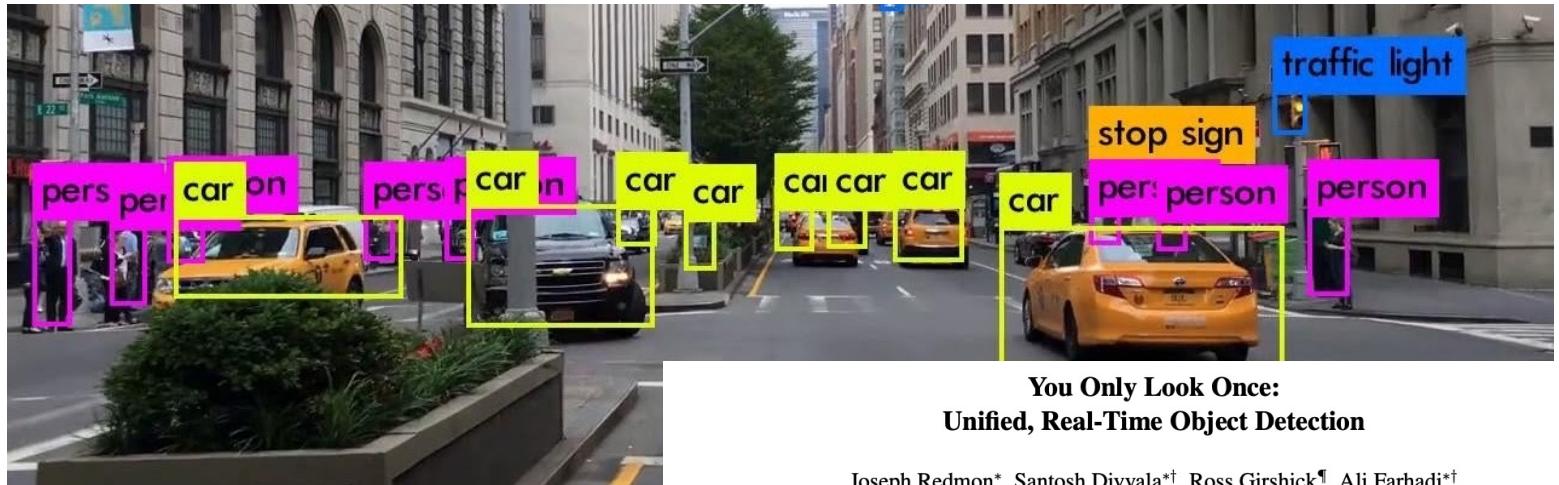
$$= \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

Faster R-CNN



Source : <https://arxiv.org/abs/1506.01497>

You Only Look Once (YOLO)



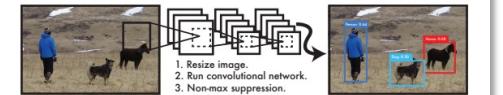
Joseph Redmon*, Santosh Divvala*†, Ross Girshick¶, Ali Farhadi*†

University of Washington*, Allen Institute for AI†, Facebook AI Research¶

<http://pjreddie.com/yolo/>

Abstract

We present YOLO, a new approach to object detection. Prior work on object detection repurposes classifiers to perform detection. Instead, we frame object detection as a re-



<https://arxiv.org/pdf/1506.02640.pdf>

$$+ \sqrt{\frac{c_e \ln(T_F)}{t_{F,a}}} \min \left(\frac{1}{4}, \frac{1}{\sum_{i=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)^2} \right)$$

Diagram illustrating a neural network architecture with layers V3, V4, V5, V6, V7, V8, V10, V11, and V12. A 'Cible' layer is shown at the top. Arrows indicate connections between layers, with some layers having multiple outgoing arrows.

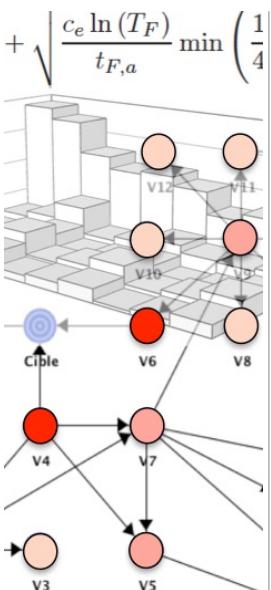
```

or (Item lCurrentItem : 
    lIdItem = lCurrentItem
    lItemConstraint = new
    for (int j = 0; j < p
        lCurrentPattern =
        lItemConstraint[j];
    }
    constraints.add(new L
        Relationship.
    )
  
```

You Only Look Once (YOLO)

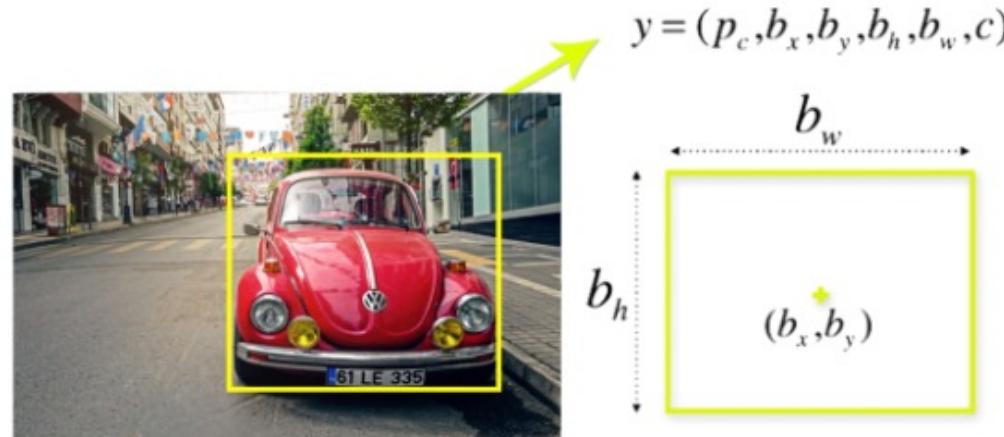
Pas de pyramide

A partir de l'image entière, YOLO renvoie des “bounding box” avec une probabilité de présence d'un objet dans cette box, ainsi que la probabilité de chaque classe d'objets possibles (80)



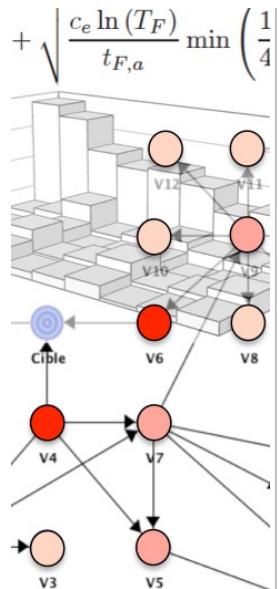
```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.
```

$$: \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$



You Only Look Once (YOLO)

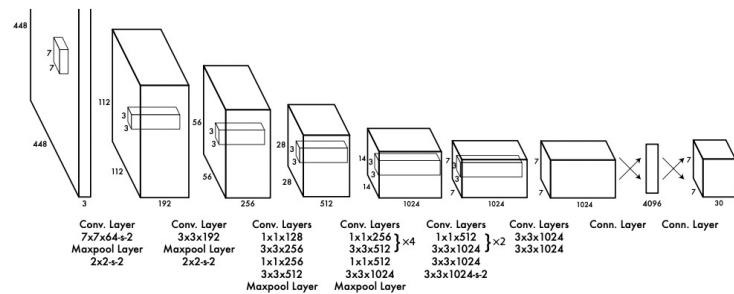
We pretrain our convolutional layers on the ImageNet 1000-class competition dataset [30]. For pretraining we use the first 20 convolutional layers from Figure 3 followed by a average-pooling layer and a fully connected layer. We train this network for approximately a week and achieve a single crop top-5 accuracy of 88% on the ImageNet 2012 validation set, comparable to the GoogLeNet models in Caffe's Model Zoo [24]. We use the Darknet framework for all training and inference [26].



```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.
```

$$= \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

Stéphane BONNEVAY – Polytech Lyon



We then convert the model to perform detection. Ren et al. show that adding both convolutional and connected layers to pretrained networks can improve performance [29]. Following their example, we add four convolutional layers and two fully connected layers with randomly initialized weights. Detection often requires fine-grained visual information so we increase the input resolution of the network from 224×224 to 448×448 .

Our final layer predicts both class probabilities and bounding box coordinates. We normalize the bounding box width and height by the image width and height so that they fall between 0 and 1. We parametrize the bounding box x and y coordinates to be offsets of a particular grid cell location so they are also bounded between 0 and 1.

We train the network for about 135 epochs on the training and validation data sets from PASCAL VOC 2007 and 2012. When testing on 2012 we also include the VOC 2007 test data for training. Throughout training we use a batch

You Only Look Once (YOLO)

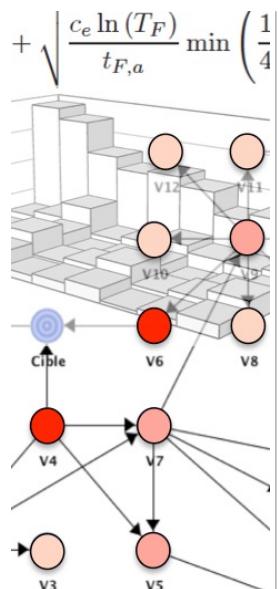
Données d'apprentissage
PASCAL VOC 2007



```

<annotation>
  <folder>VOC2007</folder>
  <filename>000037.jpg</filename>
  <source>
    <database>The VOC2007 Database</database>
    <annotation>PASCAL VOC2007</annotation>
    <image>flickr</image>
    <flickrid>338124795</flickrid>
  </source>
  <owner>
    <flickrid>seudolog</flickrid>
    <name>?</name>
  </owner>
  <size>
    <width>500</width>
    <height>375</height>
    <depth>3</depth>
  </size>
  <segmented>0</segmented>
  <object>
    <name>dog</name>
    <pose>Left</pose>
    <truncated>0</truncated>
    <difficult>0</difficult>
    <bndbox>
      <xmin>61</xmin>
      <ymin>96</ymin>
      <xmax>464</xmax>
      <ymax>339</ymax>
    </bndbox>
  </object>
</annotation>

```



```

or (Item lCurrentItem :  
  lIdItem = lCurrentItem  
  lItemConstraint = new  
  for (int j = 0; j < p  
    lCurrentPattern =  
    lItemConstraint[j];  
  }  
  constraints.add(new L  
    Relationship.  

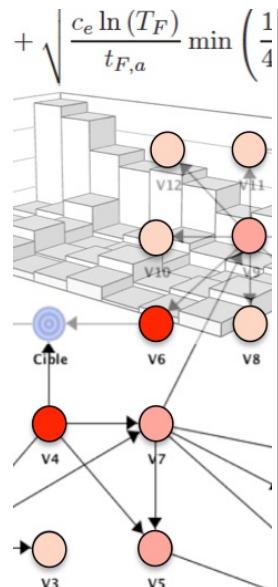

```

$$: \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

Stéphane BONNEVAY – Polytech Lyon

Deep Learning - Transfer Learning

Détection d'objets

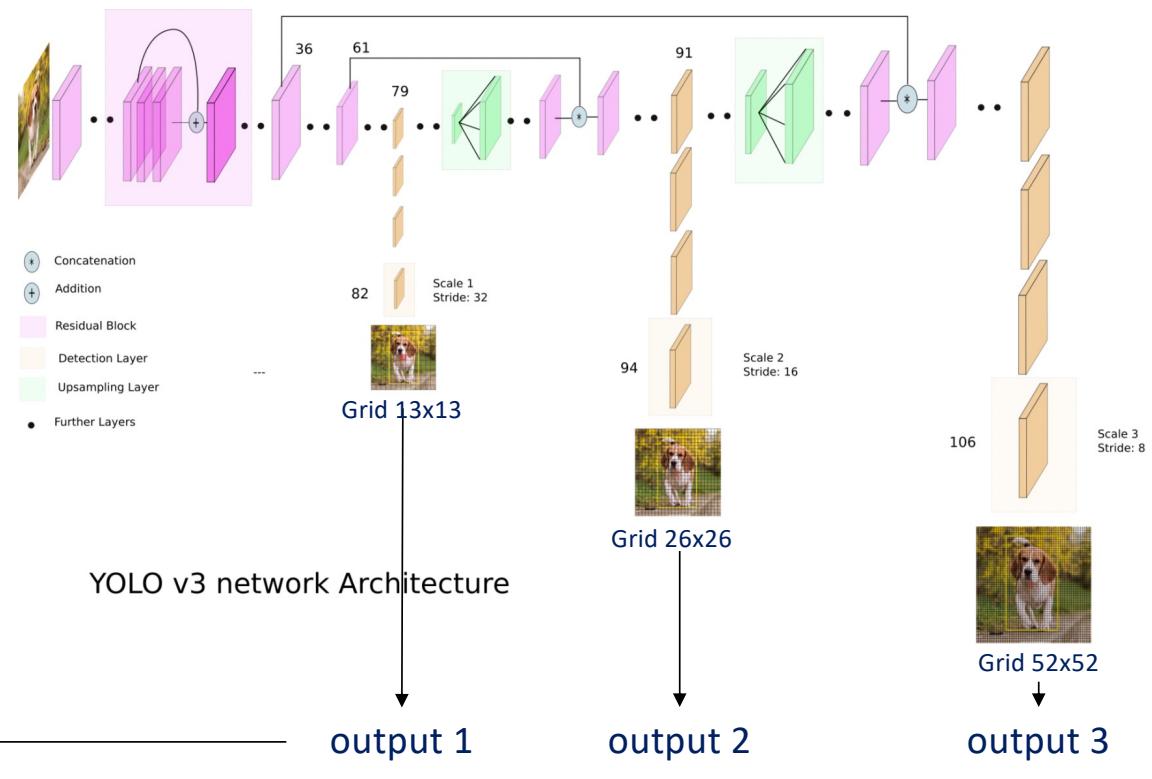


```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.  
    )
```

$$= \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

You Only Look Once (YOLO)

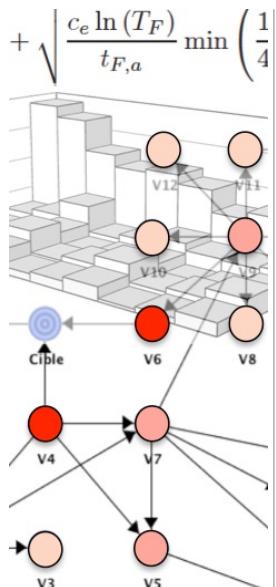
YOLO v3



3 predictions de
“bounding box” à 3
échelles différentes

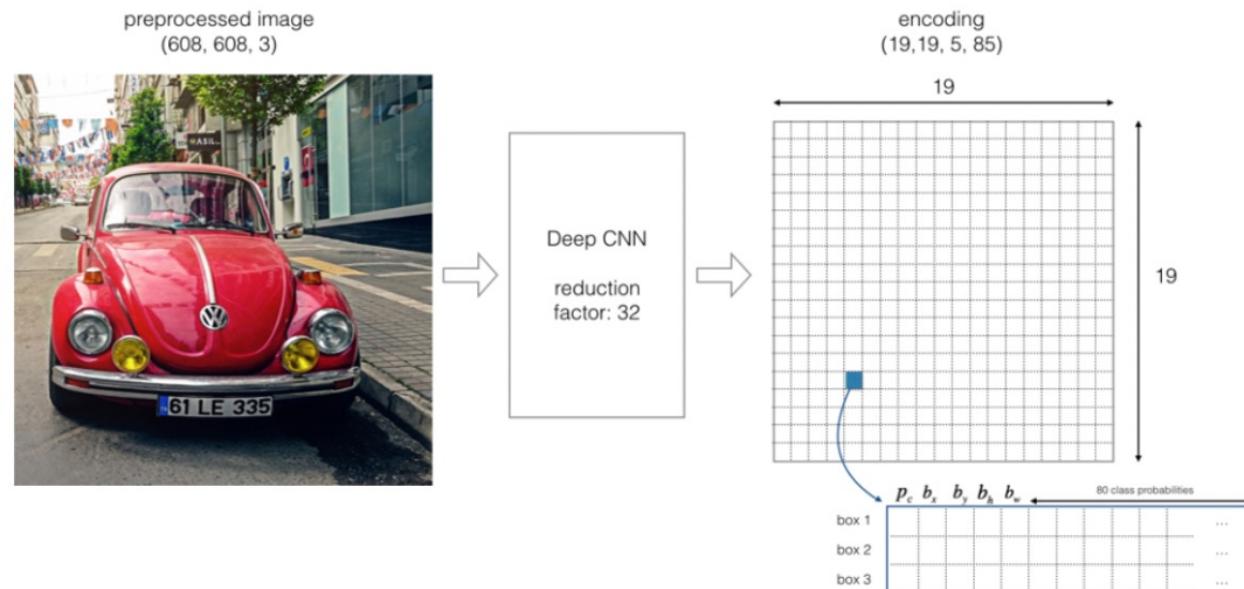
You Only Look Once (YOLO)

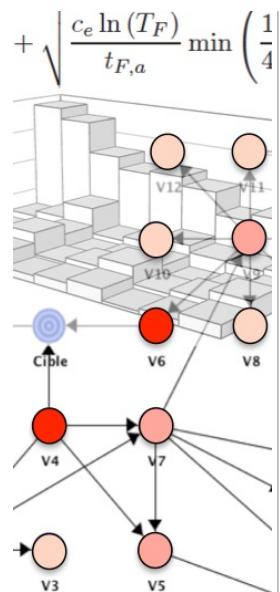
L'image est divisée en cellules qui sont chacune responsable de la prédiction de plusieurs “bounding box” (réalisé à plusieurs échelles : niveaux dans le CNN)



```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.  
    )
```

$$= \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$





```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.
```

$$= \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

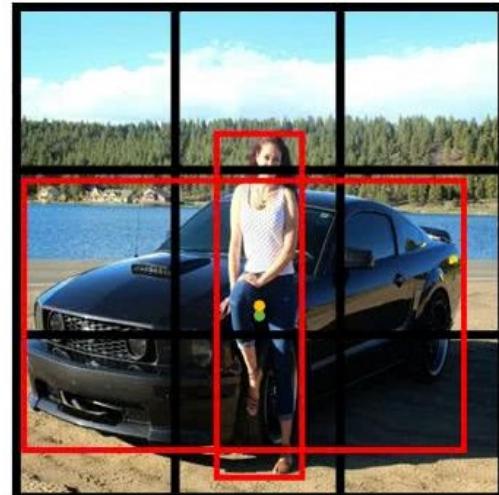
You Only Look Once (YOLO)

La voiture au centre est positionnée sur 6 cellules mais est affectée à la cellule verte, car le centre de la voiture (centre de la “bounding-box” rouge) est dans cette cellule.

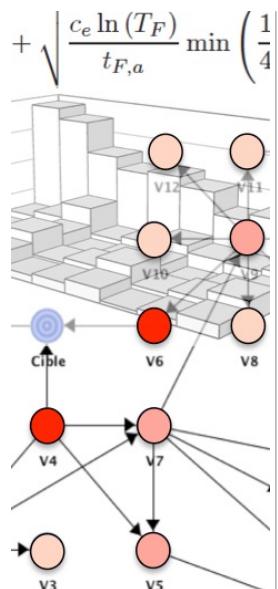


You Only Look Once (YOLO)

Une cellule peut contenir plusieurs centres de “bounding-box” et donc représenter plusieurs objets.

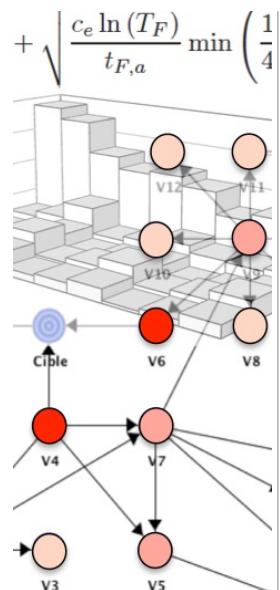


$$y = \begin{bmatrix} p_c \\ b_x \\ b_y \\ b_h \\ b_w \\ c_1 \\ c_2 \\ c_3 \\ p_c \\ b_x \\ b_y \\ b_h \\ b_w \\ c_1 \\ c_2 \\ c_3 \end{bmatrix} = \begin{bmatrix} 1 \\ b_x \\ b_y \\ b_h \\ b_w \\ 1 \\ 0 \\ 0 \\ 1 \\ b_x \\ b_y \\ b_h \\ b_w \\ 0 \\ 1 \\ 0 \end{bmatrix} \quad \left. \begin{array}{l} \text{Anchor box 1} \\ \text{Pedestrian} \end{array} \right\} \quad \left. \begin{array}{l} \text{Anchor box 2} \\ \text{Car} \end{array} \right\}$$



```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.
```

$$\prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$



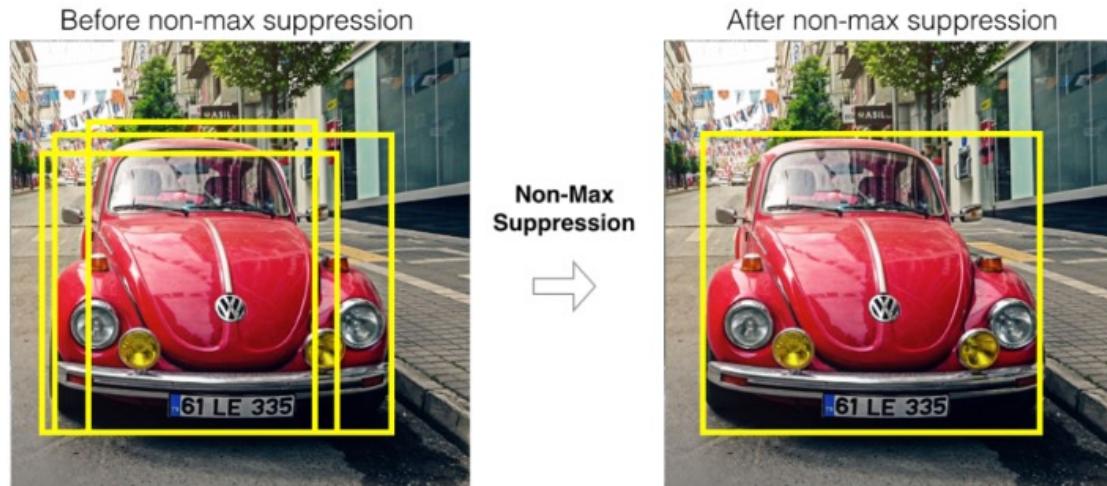
```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.
```

$$: \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

Stéphane BONNEVAY – Polytech Lyon

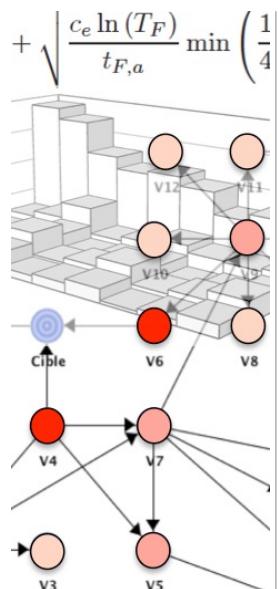
You Only Look Once (YOLO)

Lorsque plusieurs “bouding-box” identifient le même objet, on ne conserve que la “bouding-box” la plus probable par “Non-Max Suppression” :



You Only Look Once (YOLO)

Notebook python ...



```
or (Item lCurrentItem :  
    lIdItem = lCurrentItem  
    lItemConstraint = new  
    for (int j = 0; j < p  
        lCurrentPattern =  
        lItemConstraint[j];  
    }  
    constraints.add(new L  
        Relationship.  
        .
```

$$= \prod_{j=1}^{q_i} \left(\frac{(r_i - 1)!}{(N_{ij} + r_i - 1)!} \right)$$

